

Week 1 – Introduction to Secure Programming and the Application Protection Plan

Date	18.09.2026 (make-up: 30.09.2026)
Learning outcomes	LO.1, LO.5
Duration	3 hours
Prerequisites	Functions, arrays and pointers in C; <code>cd</code> and <code>ls</code> in PowerShell or a Linux terminal
Labs	code/week-01 – 4 demos; on Windows <code>.\demo.ps1</code> , on WSL / Linux <code>sh demo.sh</code> ; also opens in Visual Studio 2022 Community



By the end of this week you will be able to

1. Explain the concepts of confidentiality, integrity, and availability and the **asset–threat–vulnerability–risk** chain with your own examples.
2. Choose an **attacker model** for a piece of software; in particular, state what changes when "the device is in the attacker's hands."
3. Write the first sections (scope, architecture, assets, threats) of the **application protection plan** for a small application.
4. Derive threats from a data flow diagram using **STRIDE** and draw an **attack tree**.
5. Show and fix the dangers a program **inherits at start-up** (environment variables, memory dumps) and **secrets left in memory**.
6. Demonstrate a **buffer overflow** and a **signed/unsigned integer** error in a working example, catch them with compiler and sanitizer tools, and fix them.
7. Recognise memory management errors (leaks, double free, UAF); securely erase secrets and prevent memory from being written to disk.
8. Explain the seven layers of application protection and the ideas of **protected code partitioning** and **secure processing under encryption**; list the processes that keep a protection plan alive (version identity, change management, trade-off log).

 Lesson flow (timing plan for the instructor)


Time	Section	What happens
0:00–0:10	Introductions	Syllabus, term project, assessment, lab setup
0:10–0:35	1–4	Security concepts, attacker models, design principles (interactive lecture)
0:35–0:50	5	Overview of application protection: seven layers, the white-box attacker, the cost of protection
0:50–1:00	Break	
1:00–1:25	6–8	Protection plan, interface and asset tables, STRIDE, attack tree, DFD, risk score
1:25–1:50	9	Worked example "Vault" (seven steps) – class exercise
1:50–2:00	Break	
2:00–2:20	10–12	Demo 1 (fooling via PATH), secure start-up, Demo 2 (password left in memory)
2:20–2:40	13–16	Process memory, buffer overflows, Demo 3 , Demo 4
2:40–2:55	17–19	Memory management, protected code partitioning, secure development process
2:55–3:00	20–23	Toolbox, project kickoff, self-check

Prepare the lab beforehand

All demos run on **Windows** (Visual Studio 2022 Community or PowerShell), on **WSL**, and on **Linux**. Setup, opening the project in Visual Studio, and the safety rules: [code/README.en.md](#). In short:

Install Visual Studio 2022 Community with the **Desktop development with C++** workload. Then, in PowerShell:

```
mkdir C:\Dersler; cd C:\Dersler
git clone https://github.com/ucoruh/cen429-secure-programming.git
cd cen429-secure-programming/code
.\build.ps1
cd week-01\01-path-kandirma
.\demo.ps1
```

In Visual Studio: **File > Open > Folder** → `code` → configuration **Windows (MSVC)** → **Build > Build All**.

```
sudo apt install -y build-essential cmake gdb valgrind binutils
git clone https://github.com/ucoruh/cen429-secure-programming.git
cd cen429-secure-programming/code
./build.sh
cd week-01/01-path-kandirma
sh demo.sh
```

Code of ethics — applies every week in this course

In this course we learn about attacks **by seeing them**. Every demo runs only inside its own folder, needs no administrator rights, and changes no setting on your computer. Only try the techniques you learn **on your own computer and on the demos provided**. Trying them on someone else's system without permission is a crime.

How the course runs

- **One term project:** a C/C++ application plus a **security guide** written "as if it were going through certification"; the midterm check happens at the **week 7** demo, the final check at the **week 15** demo.
- **Two quizzes:** **week 8** (weeks 1–6 material) · **week 16** (weeks 9–14 material).
- Grade formula: $\text{Midterm} = 0.6 \cdot \text{Project1} + 0.4 \cdot \text{Quiz1}$ · $\text{Final} = 0.7 \cdot \text{Project2} + 0.3 \cdot \text{Quiz2}$.
- Every week follows the same pattern: **buggy code** → **attack** → **fix**.

0. Before we start

This section prepares you for the week. As this is the first week of the course, it builds on no earlier topic: the concept map below defines each of this week's concepts in one sentence and links it to the section where it is explained in full.

This week's concept map

Concept	In one sentence	Detail
Security	Security means protecting a system against a deliberate attacker who wants to harm it; unlike an ordinary bug, there is a conscious adversary on the other side.	§1
Asset	An asset is anything worth protecting — data, a key, a function — and security always begins with "what are we protecting?"	§1
Threat and vulnerability	A threat is a possible event or person that could harm an asset; a vulnerability is the weak point the threat can exploit; threat + vulnerability + attacker together make up risk.	§1
The CIA triad	Security's three core goals are confidentiality (only someone authorized can see it), integrity (it cannot be changed without authorization), and availability (it works when it needs to).	§1
Attacker model	An attacker model is a systematic answer to "what can the attacker see, what can they do?"; the defence is designed accordingly.	§3
White-box / MATE	MATE (Man-At-The-End) is an attacker who owns the program itself : they can read the code, see the memory, and modify it — the real situation for mobile/desktop applications.	§3
Secure design principles	The principles Saltzer and Schroeder wrote in 1975 — least privilege, secure default, open design, and others — are still the backbone of secure design today.	§4
Defence in depth	Defence in depth means building many complementary protection layers instead of relying on a single countermeasure; if one is bypassed, the next one takes over.	§5
Trade-off	Every protection carries a cost (speed, complexity, expense); security is not "infinite protection" but balanced protection with an explicitly written-down remaining risk.	§5
Threat modelling	Threat modelling is systematically answering "what are we protecting, against whom, and how?"; it lists assets, threats, and countermeasures.	§7
STRIDE	STRIDE is a method that classifies threats with six letters: spoofing, tampering, repudiation, information disclosure, denial of service, and elevation of privilege.	§7
Attack tree	An attack tree is a tree that breaks a goal down into the sub-steps an attacker could follow; the root is the attacker's goal, and the branches are the ways to reach it.	§7
Data flow diagram (DFD)	A DFD is a diagram that shows how data flows through a system via processes, data stores, external entities, flows, and trust boundaries; it is a map for finding threats.	§8

1. What is security?

Security means protecting a system against a deliberate attacker who wants to harm it; unlike an ordinary software bug, there is a **conscious adversary** on the other side. To make this concrete, think of a bank's safe. The money inside the safe is the **asset**. Someone who wants to steal the money is the **threat**. The safe's lock being faulty is a **vulnerability**. The thief using that faulty lock is the **attack**. **Risk** is the answer to the question "how likely is this event, and how much damage does it cause if it happens?"

A short history: where did the idea of secure programming come from?

- **1975** — Saltzer & Schroeder publish the **eight principles** of secure design (least privilege, defence in depth, open design...). Still foundational today.
- **1970s–80s** — the **CIA triad** (confidentiality · integrity · availability) emerges from military/government information security and becomes a common language.
- **1998–99** — the **STRIDE** threat classification (Kohnfelder & Garg) is developed at Microsoft; in the same years, Bruce Schneier introduces **attack trees**.
- **2001–2003** — Viega & McGraw's *Building Secure Software* and Viega & Messier's *Secure Programming Cookbook for C and C++* (the main source for this course) establish the culture of "build it in from the start, don't bolt it on afterward."


All of this week's tools (CIA, attacker model, STRIDE, attack tree) are products of this lineage.

Software has the same chain; only the program stands in for the safe, and data stands in for the money.

Concept	Definition	Safe example	Software example
Asset	Anything worth protecting	The money in the safe	A user's password, a payment key, customer records
Threat	A possible event or person that could harm the asset	The thief	Someone who wants to steal the password
Vulnerability	The weakness the threat can exploit	The faulty lock	A <code>strcpy</code> whose length is never checked
Attack (exploit)	Actually exploiting the vulnerability	Picking the lock	Becoming administrator with a long username
Countermeasure	A measure that reduces risk	A new lock, an alarm	Length checking, compiler protection
Risk	Probability × impact	"With this lock, in this neighborhood, there is one burglary a year"	"If this bug is found, every account is compromised"

Tehdit · zafiyet · varlık · karşı önlem

risk = olasılık × etki



Zafiyet olmadan tehdit zarar veremez; karşı önlem zafiyeti kapatır.

The three goals of security: CIA

When we talk about security, we are almost always trying to protect these three things; together they are known as the **CIA triad** (confidentiality, integrity, availability):

Goal	Question	What happens if it breaks?	Typical countermeasure
Confidentiality	Can only authorized parties see the data?	Passwords, card details leak	Encryption, access control
Integrity	Can the data be changed without authorization?	A balance, a grade, a price is changed	Hash, MAC, digital signature
Availability	Does the service work when needed?	The system crashes, nobody can use it	Redundancy, resource limiting

Three more concepts are often added to these: **authentication** (are you really who you say you are?), **authorization** (do you have permission to do this?), and **accountability / non-repudiation** (who did this, and can they deny it afterward?).

 Let's discuss in class (5 minutes)

Think about the banking app on your phone. Which CIA goal does each of the following break?

1. The app writes your account balance to the phone's memory without encrypting it.
2. Someone can change the recipient IBAN of a money transfer while it is in transit.
3. The server crashes from overload on a holiday night.

 Answers

1. Confidentiality · 2. Integrity · 3. Availability

The difference between a bug and a vulnerability

Not every software bug is a security hole; but the vast majority of security holes **are** ordinary software bugs. Once a bug lets an attacker achieve the outcome **they** want, it becomes a **vulnerability**.

 From the real world: Heartbleed (2014)

In a "heartbeat" message in the OpenSSL library, the length the client claimed was never compared against the length of the data it actually sent. When the attacker said "send me 64 KB back" while sending only a few bytes, the server sent back whatever neighboring data sat in its memory — passwords, session cookies, even private keys (CVE-2014-0160). A single missing length check affected a significant part of the internet. This week, in [Demo 4](#), we'll see a bug from the same family with our own hands.

2. Why "secure programming"?

Firewalls, antivirus software, network monitoring... All of these matter, but they **cannot rescue a badly written program**. The attacker comes in through the door that is left open — the program itself. Secure programming treats security not as a layer bolted on afterward, but as **decisions made while writing the code**:

Neden “güvenli programlama”? Beş karar, kodun içinde

güvenlik duvarı hatalı yazılmış programı kurtaramaz

Girdiye güvenme

uzunluk, tür, aralık, izin verilen karakterler — her seferinde

Sırrı kısa tut

işin bitince güvenle sil; bellekte bekletme

Ortama güvenme

ortam değişkenleri, dosya izinleri, çağrılan programlar

Güvenli tarafta kal

hata olursa reddet, kapat, sızdırma

Korumaları aç ve sına

derleyici bayrakları, sanitizer, fuzzing

Saldırgan izin verilen kapıdan girer: programın kendisinden.

- Don't trust input; **validate every input** (length, type, range, allowed characters).
- Keep secrets in memory for the **shortest time possible**, and **erase them securely** once you're done.
- Don't trust the environment the program runs in (environment variables, file permissions, the programs it invokes).
- If an error occurs, stay on the **safe side**: refuse, shut down, don't leak.
- **Turn on** the protections the compiler and operating system offer, and **test with tools** (sanitizers, fuzzing).

Over the term we will apply these principles first in C/C++, then in Java, and then in the world where "the program is in the attacker's hands" (runtime protection, code obfuscation, white-box cryptography).


3. Who is the attacker, and what can they reach?

The first question in evaluating any security measure is: "**Against whom?**" The same countermeasure may be enough against one attacker and useless against another. That's why every project starts by writing down an **attacker model**.

Attacker	What can they reach?	What can they do?	Example	Where the defence weight goes
Remote attacker	Only the messages they send over the network	Sends malformed/long input, eavesdrops on traffic	Someone attacking a web server	Input validation, TLS
Local user	An ordinary account on the same computer	Reads files, changes environment variables, runs programs	Another student on a shared server	Permissions, secure start-up
Malicious insider	Authorized access to source code, to servers	Adds a backdoor, exfiltrates data	An employee abusing their role	Privilege separation, audit logging
Attacker who owns the device (white-box, "Man-At-The-End")	Full access to the program itself	Steps through it with a debugger, dumps memory, modifies the binary	Someone trying to break a mobile payment app on their own phone	Runtime protection, code obfuscation, white-box cryptography

Saldırgan modelleri: ne görebiliyor?

soldan sağa görünürlük artar → savunma zorlaşır



Bu derste varsayılan model BEYAZ KUTU'dur: saldırgan cihazın sahibidir.

What makes this course different

Classic security courses mostly deal with the first two attackers. In this course we will give special weight to the last row — the situation where the program runs **on the attacker's own device**. Mobile payment, digital content protection (DRM), anti-cheat for games, and license enforcement all live in this world. Here, "just hide the key and you're done" is not enough: the attacker can read memory, modify the code, and stop and restart the program however they like.

4. The basic principles of secure design

The principles Saltzer and Schroeder wrote in 1975 are still the backbone of secure design today. Keep each one in mind with a one-sentence example:

Saltzer ve Schroeder (1975): güvenli tasarımın omurgası

elli yıl sonra hâlâ geçerli sekiz ilke

En az ayrıcalık
işini yapacak kadar yetki

Güvenli varsayılan
varsayılan: izin YOK

Tam aracılık
her erişim her seferinde denetlenir

Açık tasarım
gizlilik anahtarda, tasarımda değil

Mekanizma ekonomisi
basit tasarım = az hata

Ortak mekanizmayı azalt
paylaşılan yol = paylaşılan risk

Yetki ayrımı
tek onay yetmez

Kullanılabilirlik
zor gelen güvenlik atlanır

Bu ilkeler bir denetim listesi değil, tasarım kararlarını tartarken kullanılan ölçülerdir.

Principle	Meaning	Example
Least privilege	Every component runs with just enough privilege to do its job	The reporting program runs as an ordinary user, not as administrator
Secure default	The default state is "no permission"	A new user's <code>yonetici</code> (admin) field starts at 0
Complete mediation	Every access is checked, every time	Permission is not checked once and then cached
Open design	Security rests on the secrecy of the key, not the secrecy of the design (Kerckhoffs)	You do not write your own encryption algorithm
Economy of mechanism	A simple design contains fewer bugs	A clear 20-line check beats a "clever" 500-line one
Separation of privilege	A critical operation requires more than one condition	Both a password and an SMS code for a money transfer
Least common mechanism	Resources shared between users are minimised	Every session gets its own temporary file
Psychological acceptability	Security must be usable, or it gets bypassed	Asking for a password every 5 minutes leads to the password being written on paper

Let's add one modern principle to these: **defence in depth**. We don't rely on a single countermeasure; we build layers that complement each other. If one is bypassed, the next one takes over. [In week 3](#) we will see how a sensitive key is carried inside exactly **four separate protection layers** (a security shell).

5. Overview of application protection: the defence-in-depth map

So far we have talked about **what** security is and **against whom** we are protecting. Now let's map out the toolbox we will use for the rest of the term. "Application protection" is not a single technique; it is a whole of complementary **layers**. Each layer answers a different question and stands against a different attacker capability.

Why a single countermeasure is not enough

Think of a building. You've fitted a good lock on the door. But if the window is open, the key is under the doormat, or the safe inside has no combination, the lock alone means nothing. It's the same in software:

- Your code can be **bug-free**, but the attacker can open the binary and **read your logic**.
- You can **hide your logic**, but the attacker can stop the running program with a debugger and **read its memory**.
- You can protect memory **against being monitored**, but if the key sits in the code as **plain text**, a single string scan is enough.
- You can hide the key, but the attacker can **change a check in the binary** (e.g. making the "is the license valid?" question always answer "yes") and bypass the protection entirely.


Each line shows what one countermeasure **does not** solve. This is the logic of defence in depth: each layer closes the path the previous one left open; the attacker has to defeat **all of them at once**, and that multiplies the cost.

Seven layers

The diagram below orders the layers we will see this term from the **inside out**. The inner layers deal with "write the code correctly"; the outer layers deal with "what happens once the code is in the attacker's hands?"

Yedi koruma katmanı

aşağıdan yukarı: her katman bir öncekinin üzerine kurulur



Bir katman aşılsa bile diğerleri durdurur — derinlemesine savunma.

Layer	Question it answers	Typical techniques	Attacker it stops
1. Secure design	What are we protecting, against whom?	Threat model, asset and interface tables, least privilege	Anyone looking for a design flaw
2. Secure coding	Is my code bug-free?	Input validation, bounds checking, safe APIs, SEI CERT rules	An attacker sending input remotely
3. Compiler/OS protections	If a bug slips through, is it harder to exploit?	Stack canary, ASLR, DEP/NX, CFI, <code>_FORTIFY_SOURCE</code>	An attacker trying to exploit a memory error
4. Obfuscation	Can someone reading the code understand it?	Control-flow flattening, opaque predicates, string obfuscation, virtualisation	A reverse-engineering analyst
5. RASP	Does the program notice it is under attack?	Debugger/hook detection, root/emulator detection, integrity checks	An attacker inspecting the program while it runs
6. Cryptography	Even if the data is captured, is it meaningless?	Authenticated encryption, key hierarchy, white-box crypto	An attacker who obtains stored or transmitted data
7. Assurance	How do we prove all of this actually works?	Code review, fuzzing, penetration testing, ETSI/EMV/FIPS/Common Criteria	Auditors and evaluators

The idea that ties the layers together: the security shell

In the method the instructor has used in the field, a sensitive piece of data (e.g. a key) is carried inside **several of these layers at once**: it is stored encrypted, decrypted only at the moment it is needed and only as much as needed, the code section where it is decrypted is obfuscated and integrity-checked, and the process it runs in has debugger detection. In this course we will call these nested layers a **security shell**. In [week 3](#) we will build a four-shell design step by step.

Which layer matters when? Choosing by attacker model

Not every application needs all seven layers equally. The choice is driven by the **attacker model** we defined in section 3:

Application type	Where is the attacker?	Priority layers
Server-side web service	On the network; no access to code or server	1, 2, 3, 6 (TLS), 7
Desktop business application	On the user's computer; but usually the user is not the attacker	1, 2, 3, 6
Mobile banking / payment app	The device may be in the attacker's hands	All of them; especially 4, 5, 6
Game, licensed software	The user themselves is the attacker (cheating, pirated copies)	Mostly 4, 5
Embedded device (e.g. a smart meter)	Physical access is possible	2, 3, 6, hardware-backed protection

The critical distinction here is: on the **server side**, your code runs on a machine **you control**; the attacker only controls the inputs. On the **client side** (phone, desktop, game console), the attacker **owns the program itself**: they can copy it, open it, modify it, and re-run it as many times as they like. This model is called the "**white-box**" attacker, and layers 4, 5, and 6 exist precisely for this attacker.

Obfuscation is not a substitute for correct code

A common mistake is to skip secure coding and say "we'll obfuscate the code anyway." Obfuscation **slows the attacker down**, it does not **fix the bug**. An obfuscated buffer overflow is still a buffer overflow; it just takes a bit longer to find. The outer layers only make sense once the inner layers are solid.

The white-box attacker's six paths

The security guide for a sensitive library running on the client side typically summarizes the attacker's goals in a six-row **attack table**. Each row is a question that at least one of the layers has to answer:

#	Attacker's path	Example	Layers that answer
1	Bypassing platform controls	A rooted device, an emulator, a modified OS	5 (RASP)
2	Reverse-engineering the code	Reading logic and keys with a decompiler	4 (obfuscation), 6
3	Modifying the code	Flipping an <code>if</code> check and repackaging	5 (integrity), 4
4	Abusing the interfaces between components	Calling the library from your own program, reading data as it crosses the interface	1 (design), 5, 6
5	Extracting assets at runtime	Memory dump, debugger, hook	2 (erasure), 5, 6 (white-box)
6	Redirecting control flow at runtime	Skipping a branch with a debugger, changing a return value	5 (flow counter), 4

This table is the **starting point** for the threats section of a protection plan: for every asset, you write down how each of these six paths is closed off.

The same set of requirements has two more items that state plainly why defence in depth is mandatory:

- The application must **not rely solely** on security functions that have no assurance of their own: features like the OS's protections, a keychain, an app sandbox, or a hardware-backed key store vary from device to device and can be disabled by the attacker.
- Mechanisms the platform provides (e.g. the OS's crypto library) can be used, but only **in addition to** the application's own protections.

Defensive programming is listed item by item in the same place: every input and output is checked; assets are processed for the shortest time possible; assets are never passed in the clear even across interfaces inside the code; delays are added against timing-correlation attacks; and a secure erase function that wipes the data is used once the work is done. We will apply every one of these rules over the course of the term.

The cost of protection: every layer has a price

Adding a layer is never free; security involves a constant **trade-off** between speed and complexity — the goal is not infinite protection but **balanced** protection, with the remaining risk written down explicitly. When writing a protection plan, you weigh these four costs for every layer:

1. **Performance:** Obfuscated code runs slower (e.g. control-flow flattening can slow a function down several times over). Integrity checking and debugger detection also cost CPU time.
2. **Size:** Obfuscation, virtualisation, and table-based white-box cryptography make the binary bigger.
3. **Development and maintenance:** Bugs in obfuscated code are hard to debug; crash reports become unreadable. That's why obfuscation is applied **only to sensitive sections**, and symbol-mapping files are kept somewhere safe.
4. **False positives:** An overly sensitive root or emulator check can punish a legitimate user (e.g. a phone with developer options turned on).

That's why every countermeasure in a protection plan is written **with its justification**: "We protect this asset, against this attacker, at this cost." A countermeasure with no written justification is either unnecessary or in the wrong place.

The path we'll follow through the semester

We will fill in this map by opening up one layer each week:

Week	Layer	Main question
1	1, 2	Secure design, threat model, how memory errors happen
2	1	Malware, security models, classifying vulnerabilities
3	6	How do we protect data in transit, at rest, and in use?
4	2, 3, 4	How do we harden C/C++ code?
5	2, 4	Injection, obfuscation, and dependency security in Java and interpreted languages
6	5	How does the program protect itself while under attack?
9	4	Advanced obfuscation and measuring its effectiveness
10	6	AES, RSA, digital signatures, PKI
11	6	How is a key protected while it is inside code in the attacker's hands?
12–13	7	Certification, penetration testing, security requirements
14	4	Automated obfuscation and diversification with Tigress



How does an evaluator look at these layers?

An independent evaluator reading a product reads this table **backward**: they try the easy paths first (string scanning, opening the binary, attaching a debugger) and note which layer stopped them. Every step that was **not** stopped becomes a finding in the report. Your project's security guide will likewise answer, for every layer, "which countermeasure did I take, and how did I test it?"

Try it yourself

1. What is the attacker model for an **offline password manager** (all data on the user's computer)? Which of the seven layers matter most? Why?
2. Under what conditions is the statement "our app runs on the server, we don't need obfuscation" true, and under what conditions is it false?
3. A game company uses **only obfuscation** to stop cheating. Which attack path stays open?



Answer hints




1. The attacker could be someone copying the file (a stolen laptop) or someone who infected the computer with malware; cryptography (6: password-based key derivation, authenticated encryption) and secure coding (2) matter most; since the user themselves is not the attacker, obfuscation is secondary.
2. It is true if the code really runs only on the server and no logic is ever sent to the client. It is false if the client application also carries business logic or a key (a mobile app, JavaScript in the browser).
3. A cheating client sending **inconsistent data** to the server: obfuscation protects the client, but if the server does not validate every value coming from the client (e.g. "this player can't move 100 meters in one second"), cheating is still possible. The rule: **the server never trusts the client**.

6. The application protection plan

The security of a piece of software is not left to chance; it is managed with a **written plan**. In products that go through security certification at an independent lab, this plan turns into a security guide hundreds of pages long. In this course, you too will write a small version of one for your term project. The backbone of the plan is:

Uygulama koruma planı: yedi adım

1–3 TANIMLA · 4–7 KORU ve KANITLA



Bu yedi adım dönem projesinin de iskeletidir (S2–S17).

The plan is not written once and forgotten: every time the product changes (a new feature, a bug fix, a new release) you return to step 1 and update the plan.

1. **Scope:** What are we evaluating? Saying "version 2.1" alone is not enough — which binary, which source code, which hash value? In certification this is called the **target of evaluation**.
2. **Architecture and interfaces:** What are the components, which channel do they talk to each other over, and how is identity verified on each channel?
3. **Assets:** every piece of data to be protected is listed one by one: where does it live, when is it created, when is it erased, does it need confidentiality (C) or integrity (I)?
4. **Threats:** who is the attacker (see section 3), and by what path can they reach each asset?
5. **Countermeasures:** for each threat, which countermeasure, in which layer?
6. **Verification:** how is it proven that each countermeasure actually works — a unit test, a bypass attempt?
7. **Residual risk:** the risks that cannot be closed, or are knowingly accepted, and why (e.g. performance).


The plan's two core tools: the interface table and the asset table

Steps 2 and 3 of the plan are written with two tables; we will use this same method throughout the whole term. Let's see the method on an example architecture: a mobile app that makes contactless payments over the phone's NFC, and a **security library** that does its critical work. Part of the library is in Java, while the most security-sensitive part runs in a **native layer written in C/C++**.

Mobil ödeme: mimari ve arayüzler

kırmızı kutu = güvenilemez ortam (saldırgan burada)

Kullanıcının telefonu — GÜVENİLMEZ



Her arayüz (A–E) ayrı bir saldırı yüzeyidir: her biri için önlem ve kanıt yazılır.

Notice two things in this diagram:

- The phone box is labelled "**untrusted environment**." The phone is in the user's — i.e., a potential attacker's — hands. It may be rooted, it may have a debugger attached. That's why the design starts from the assumption "**I do not trust the phone**."
- Every arrow labelled with a letter is an **interface**. The following table is filled in for each interface:

Interface	End A	End B	Authentication	Confidentiality / integrity
A	Mobile app	Library (Java)	The calling app's signature is verified	Same process; no sensitive data crosses
B	Library (Java)	Library (native)	Mutual integrity check	Data crosses encrypted and masked
C	Library (native)	Local database	—	Encryption + integrity code (MAC), key bound to the device
D	Library	Back-end server	Certificate pinning + server verification code	TLS and message-level encryption on top
E	Library	Payment terminal	Payment protocol	One-time payment key

And a slice of the **asset table** (the values are synthetic):

Asset	Type	Where	Creation → deletion	Protection
One-time payment key	Cryptographic, dynamic	Database (encrypted), memory at time of use	Downloaded from the server → used in one payment → erased	C, I
Device fingerprint	Dynamic	Memory, in two separate pieces	Computed at app start-up	I
Session key	Cryptographic, dynamic	Memory only	Derived when the session with the server is set up → erased at session end	C, I
Transaction counter	Dynamic	Database	Increments on every payment	I
Server public key	Cryptographic, static	Embedded in the app	At build time	I

Why so much detail?

The first thing an evaluator does is look at this list and ask, for every row: "**Where and when** is this asset exposed?" An asset that is not on the list is an asset that is not protected. In your project, too, every asset will have a place in this table.



What comes after this method

Every week this term we will open up one of the methods that protect the assets in these tables: layered data protection and security shells ([week 3](#)), hardening native code (weeks 4 and 9), hardening the Java side ([week 5](#)), runtime self-protection ([week 6](#)), crypto methods and key hierarchy ([week 10](#)), white-box cryptography ([week 11](#)), certification and requirement mapping (weeks 12–13).

7. Threat modelling: STRIDE and attack trees

Threat modelling means asking, **systematically**, "if I were the attacker, how would I attack this?" One of the most common methods is **STRIDE**, developed at Microsoft: each letter is a family of threats and the security property it breaks.

Letter	Threat	Property it breaks	Example (mobile payment architecture)
S	Spoofing — identity impersonation	Authentication	A fake app disguising itself as the real app and calling the security library
T	Tampering	Integrity	A check in the native library being altered in the binary
R	Repudiation	Non-repudiation	A user claiming "I didn't make this payment" with no proof to the contrary
I	Information disclosure	Confidentiality	The payment key being read from a memory dump
D	Denial of service	Availability	Thousands of fake registration requests against the server
E	Elevation of privilege	Authorization	An ordinary user reaching an administrator function


Applying it is simple: draw the system's **data flow diagram** (external entities, processes, data stores, data flows, and **trust boundaries**), then ask all six letters for every element, one by one. Every arrow that crosses a trust boundary is where the most threats come from.

Attack tree

In an attack tree, the root is the attacker's **goal**; the branches below it are the ways to reach that goal. In an **OR** node one branch is enough; in an **AND** node all of them are required. Once the tree is drawn, the cheapest path and the key points to defend fall out on their own.

Saldırı ağacı: ödeme anahtarını ele geçir

VEYA = tek dal yeter · VE = tüm dallar gerekir



Two conclusions follow immediately from this tree: (1) Path C is expensive because it requires breaking two layers at once – that's defence in depth at work. (2) Path B has three alternatives; that means the key **in memory** is the weakest point, and it needs dedicated countermeasures (this week's Demo 2, runtime protection in [week 6](#), white-box in [week 11](#)).

Class exercise (15 minutes, groups of 3–4)

Think of a "student grading system": the instructor enters grades, the student sees their grade, the data lives on a server.

1. Draw the data flow diagram; mark the trust boundaries with a dashed line.
2. Write at least six STRIDE threats (one for each letter).
3. Draw an attack tree for the goal "raise my grade from 45 to 85." Which is the cheapest path?

Hint


Trust boundaries: browser ↔ server, server ↔ database, student role ↔ instructor role. For "E," think of a student reaching the grade-entry page by typing its address directly; for "R," think of there being no record of "who changed this grade, and when?"

8. Going deeper into threat modelling: the data flow diagram, STRIDE in detail, and risk scoring

In the previous section we met STRIDE and the attack tree. In this section we'll learn to apply the method **step by step, like an engineer**: how do we draw the diagram, which letters do we ask of each element, and how do we rank the dozens of threats that come out of it?

Veri akış diyagramı: dört öge ve güven sınırı

tehdit modellemenin çizim dili



Sınırı geçen her ok bir tehdit adaydır: orada STRIDE'in altı harfini tek tek sorun.

Kural: güven sınırını çizmeden tehdit listesi çıkarılmaz.

The four elements of a data flow diagram (DFD) and the trust boundary

The diagram used in threat modelling is a simplified version of the data flow diagram you know from software engineering courses. There are only five symbols:

Symbol	Name	What it represents	Example
Rectangle	External entity	A person or system not under your control	User, payment terminal, email server
Circle / rounded box	Process	Code that processes data	Login module, encryption library
Two parallel lines	Data store	Where data sits	Database, config file, registry
Arrow	Data flow	Data moving from one place to another	HTTP request, function call, file write
Dashed line	Trust boundary	Where the level of trust changes	Internet ↔ server, user mode ↔ kernel, application ↔ operating system

The concept of the **trust boundary** is the heart of the method. The moment a piece of data crosses a trust boundary, it becomes "data that came from the other side" and cannot be used without validation. Most threats are born at these crossings. A practical rule: color every arrow that crosses a trust boundary red in your diagram, and start your review there.

Which letter applies to which element?

Not every element is exposed to all six threats. The "STRIDE-per-element" table that Microsoft popularized shows which questions actually make sense:

Element	S	T	R	I	D	E
External entity	✓		✓			
Process	✓	✓	✓	✓	✓	✓
Data store		✓	✓ ¹	✓	✓	
Data flow		✓		✓	✓	

¹ If the data store is an audit log, the repudiation threat is asked too: if the record is deleted or altered, who did what can no longer be proven.

The logic of the table is intuitive: a data flow (e.g. a network packet) cannot be **spoofed**, because it has no "identity"; but it **can** be eavesdropped on (I), altered (T), or cut off (D). A process, on the other hand, is exposed to all six threats: it can be impersonated, its code can be altered, its memory can be read, it can be crashed, its privilege can be abused.

Questions to ask for each letter, and typical countermeasures

Letter	Ask yourself	Typical countermeasure	Where it appears in this course
S	Is the other party really who they claim to be? Could someone else pose as them?	Authentication, mutual TLS, signature verification, certificate pinning	Weeks 3, 10
T	Can this data or code be altered in transit or on disk? Would I notice if it were?	MAC/HMAC, digital signature, integrity checking, read-only memory	Weeks 2, 3, 6
R	If a user says "I didn't do this," do I have proof?	Tamper-resistant audit log, signed transaction, timestamp	Week 2
I	Can this data end up in the hands of someone unauthorized — on disk, on the network, in memory, in a log, in an error message?	Encryption, memory wiping, data minimization, masking	Weeks 1, 3, 11
D	Can this component be crashed, deadlocked, or have its resources exhausted?	Input limits, rate limiting, timeouts, resource quotas	Weeks 1, 4
E	Can someone less privileged get someone more privileged's work done?	Least privilege, complete mediation, memory safety, sandboxing	Weeks 1, 2, 4

The same bug can have more than one letter

A buffer overflow is a threat under **D** (the program crashes), **E** (the attacker hijacks control flow and runs privileged code), **and I** (a secret in adjacent memory is read) all at once. STRIDE is not a classification, it's a **thinking tool**: its purpose is not to put a threat in the right box, but to make sure **no threat is skipped**.

Ranking threats: likelihood × impact

Threat modelling on a medium-sized system easily produces 50–100 threats. You cannot close all of them at once. That's why every threat gets a **risk score**. The simplest method is to score likelihood and impact from 1 to 3 and multiply them:

	Impact: Low (1)	Impact: Medium (2)	Impact: High (3)
Likelihood: High (3)	3 — medium	6 — high	9 — critical
Likelihood: Medium (2)	2 — low	4 — medium	6 — high
Likelihood: Low (1)	1 — low	2 — low	3 — medium

When determining **likelihood**, ask what the attacker needs: does it require physical access, or can it be done over the internet? Does it need special tooling or expertise? Can the attack be automated? When determining **impact**, go back to

the asset table: which asset is affected, how many users are affected, what is the financial or legal consequence?

More detailed scales

The industry uses more detailed scoring systems than this. **CVSS**, which we'll see in week 2, scores a vulnerability from 0–10 using criteria such as attack vector, complexity, privileges required, user interaction, and impact on confidentiality/integrity/availability. In payment system certification, a scale called "**attack potential**" is used, which adds up the time required, expertise, knowledge of the target, window of opportunity, and equipment ([week 13](#)). This week, the simple 3×3 matrix is enough.

Four responses to a threat

For every threat, you consciously choose one of four options and write it into the plan:

1. **Mitigate:** Add a countermeasure. The most commonly chosen response.
2. **Eliminate:** Remove entirely the feature that causes the threat. E.g. if a "remember my password" feature isn't necessary, never writing it at all zeroes out the threat of the password being stored on disk. This is the strongest response.
3. **Transfer:** Hand the risk off to someone else. E.g. never storing card data at all and leaving it to the payment service provider, or an insurance policy.
4. **Accept:** Knowingly take on the risk and write down the justification. E.g. "We take no countermeasure against the chip being physically opened and examined; the cost of this attack is higher than the value it protects."

An accepted risk is not a forgotten risk

Saying "accept" means writing the risk into the plan's **residual risk** section and re-evaluating it in the next release. A risk that is not written into the plan has not been accepted — it has been **overlooked**. Evaluators look first at whether a residual risk list exists at all, and at its justifications.

9. Worked example: writing a protection plan step by step

Now let's apply all the tools from the previous three sections, start to finish, on a small application. This example will serve as a **template** for the first deliverable of your term project.

Step 0 — Define the application

Let "**Vault**" be a password vault, written in C++, running on the desktop:

- The user sets a **master password**. The vault encrypts the records inside it (site name, username, password) with a key derived from this master password and stores them in a **single file**.
- The user can copy a record's password to the **clipboard**.
- Optionally, the encrypted vault file can be sent to a **backup server**.
- The application downloads new versions from an **update server**.

Step 1 – Scope


Field	Value
Target of evaluation	kasa.exe / kasa version 1.0, core library kasa_cekirdek, configuration file
Platforms	Windows 10/11 (x64), Ubuntu 22.04 (x64)
Out of scope	The backup server's own security (owned by a separate team), the operating system itself
Security objective	The vault file and its backups must be unreadable and undetectably unmodifiable to anyone who does not know the master password

The "out of scope" row matters: writing down explicitly what you are **not** protecting is as valuable as writing down what you are.

Step 2 – Architecture and interfaces

Parola kasası: bileşenler ve arayüzler

her ok bir arayüz — her arayüz doğrulanır



Interface	End A	End B	Authentication	Confidentiality / integrity
A	User	Interface	Master password	Password is never shown on screen
B	Interface	Core	Same process	Master password is erased from memory after the call
C	Core	Vault file	–	Authenticated encryption (AES-GCM), key derived from the master password
D	Core	Clipboard	–	Clipboard is cleared after 30 s
E	Core	Backup server	Server certificate verified, user token	TLS + the file is already encrypted
F	Interface	Update server	The update package's digital signature is verified	TLS; the real assurance is the signature

Step 3 – Assets

#	Asset	Type	Where	Creation → deletion	Protection
V1	Master password	Secret, short-lived	Memory only	User types it → key is derived → erased immediately	C
V2	Vault key	Cryptographic, session-lived	Memory only	Derived → erased when the vault locks	C, I
V3	Vault records (decrypted)	Secret	Memory (only the record being displayed)	Opened → erased once off screen	C
V4	Vault file	Encrypted data	Disk, backup server	Persistent	C, I
V5	Salt and derivation parameters	Public, but integrity matters	Vault file header	When the vault is created	I
V6	Update signature public key	Cryptographic, static	Embedded in the app	At build time	I
V7	Backup token	Credential	Disk (OS key store)	When the user connects → until revoked	C

The "Protection" column uses four values: **C** (confidentiality), **I** (integrity), **I+** (integrity plus accountability and authentication: the data's **origin** must also be verified; in the field this is provided by an authentication code shared with the server), and **N** (no protection needed). Asset tables used in the field also record each asset's size, the source of its value, its default value, and which table or file it lives in; static, dynamic, and cryptographic assets are listed in separate tables.

Step 4 – Threats (STRIDE, interface by interface)

#	Where	Letter	Threat	Likelihood	Impact	Risk
T1	C / V4	I	The vault file is copied from a stolen laptop, a weak master password is brute-forced offline	3	3	9
T2	B / V1, V2	I	The master password or key stays in memory; it's read from a memory dump, a swap file, or a hibernation file	2	3	6
T3	D / V3	I	The password copied to the clipboard is read by another application	3	2	6
T4	C / V4	T	The vault file is modified; the app uses the corrupted record without noticing, or crashes	2	2	4
T5	C / V5	T	The derivation parameter (iteration count) in the header is dropped to 1; the next save uses a weak key	1	3	3
T6	F / V6	T, E	A fake update package is	2	3	6

#	Where	Letter	Threat	Likelihood	Impact	Risk
			installed; the attacker's code runs with the user's privileges			
T7	E	S	A fake backup server steals the token	2	2	4
T8	Process	T	A memory error in the vault file parser; a malicious file crashes the program or hijacks it	2	3	6
T9	A	S	Someone next to the user reads the master password off the screen (shoulder surfing)	1	3	3

Step 5 — Countermeasures

Threat	Response	Countermeasure	Layer
T1	Mitigate	Derive the key from the master password with a slow derivation function (PBKDF2 with a high iteration count, or Argon2id), a minimum password length requirement, a password strength meter	6, 2
T2	Mitigate	Securely erase the key and password right after use, lock the memory against being swapped, disable crash dumps (sections 11–12 and 17 of this week)	2
T3	Mitigate	Clear the clipboard after 30 s; use the "exclude from clipboard history" flag if the platform supports it	1
T4	Mitigate	Authenticated encryption: never open the file if the tag doesn't match	6
T5	Mitigate	Make the header additional authenticated data (AAD) of the encryption; hard-code a minimum iteration count and reject anything below it	6, 2
T6	Mitigate	Never run the update package without verifying its signature against the embedded public key; reject version downgrades	6
T7	Mitigate	Certificate validation + hostname check, optional certificate pinning	6
T8	Mitigate	Validate every length field in the parser, test with fuzzing, turn on compiler protections	2, 3, 7
T9	Accept	Mask the password on screen; further countermeasures against physical observation are out of scope	—

Step 6 — Verification plan

For every countermeasure, you write down the answer to "how do I prove this works?" If a countermeasure has no test, it **counts as if it doesn't exist** in the evaluator's eyes.

Countermeasure	Test	Expected result
T1 derivation	Unit test: derivation time is measured	≥ 250 ms on the target machine
T2 erasure	After the vault locks, the process memory is dumped and searched for the master password	Not found
T4 integrity	A random byte of the vault file is modified	The app says "file corrupted," refuses to open, does not crash
T5 parameter	The iteration count in the header is set to 1	The app rejects the file
T6 signature	A package with a broken signature is presented	Installation is refused, the event is logged
T8 parser	24 hours of fuzzing (week 4)	No crashes; no sanitizer findings

Step 7 — Residual risk

Risk	Why was it accepted?	Re-evaluation
Memory can be read if malware is running on the computer while the vault is open	Full protection against malware running with the same user privilege is not possible; the auto-lock timeout was kept short	Every release
Shoulder surfing (T9)	The physical environment is out of scope	—
A very weak master password	The user is warned but not forcibly blocked (usability)	Based on user feedback

✓ Carry the template into your project

This seven-step structure is the **skeleton** of your term project's security guide. This week you are only expected to fill in **steps 0–4** for your own project (see "Term project: this week"). You will fill in steps 5–7 throughout the term with the countermeasures you learn each week.

Common mistakes

- **Forgetting assets:** writing only "user data" and never listing keys, tokens, counters, configuration files, and **the code itself** (which is an asset whose integrity must be protected).
- **Drawing the trust boundary wrong:** putting the client application you wrote yourself in the "trusted" box. If the client runs on the user's machine, it is **untrusted**.
- **Not matching a countermeasure to a threat:** saying "we use TLS" is incomplete without saying which threat TLS closes (eavesdropping and tampering) and which it does **not** (data at the server, the client itself).
- **A countermeasure with no test:** saying "we erase the password from memory" and never checking whether the compiler removed that erasure (this week's Demo 2 is exactly this).

10. At program start-up: an untrusted environment

When a program starts running, it does not start with a **blank page**. It inherits a lot from whoever launched it: **environment variables** (`PATH`, `LD_PRELOAD`, `IFS` ...), **open file descriptors**, the working directory, `umask`, resource limits. If the program trusts these, then whoever controls them also controls the program's behaviour. The first chapter of our source book (Viega & Messier, recipes 1.1–1.9) deals with exactly this topic of "secure start-up."

Demo 1 – Fooling the program via PATH


i Demo 1 · `code/week-01/01-path-kandırma` · CWE-426 (untrusted search path)

The program calls a system command: `date` on Linux, which prints the date, and `hostname` on Windows, which prints the computer's name. The question is: **which** `date`, **which** `hostname` ?

When you call a command without an absolute path, the shell searches the folders in the `PATH` variable **from left to right** and runs the **first one it finds**:

PATH ile kandırma: neden system() tehlikeli?

kabuk, programı adına göre PATH sırasıyla arar



On Windows the situation is even worse: `cmd.exe` searches the **working directory before** `PATH`, and it also tries extensions like `.bat` and `.cmd`. Even **which folder** you run the program from can change **which file** actually runs.

Buggy code:

rapor.c

```

#ifdef _WIN32
#define KOMUT "hostname"
#else
#define KOMUT "date"
#endif

int main(void)
{
    printf("=== Aylik Satis Raporu ===\n");
    fflush(stdout);

    /* HATA: çıplak komut adı + kabuk + devralınan ortam (PATH, ...) */
    int durum = system(KOMUT);
    ...
}
  
```

Run the demo:

```
cd code\week-01\01-path-kandirma
.\demo.ps1

cd code/week-01/01-path-kandirma
sh demo.sh
```

Output – WSL (abridged; on Windows the same flow with hostname)

```
ADIM 1 – Normal calisma: program gercek date'i buluyor
=== Aylik Satis Raporu ===
Rapor tarihi: Sat Sep 19 19:11:32 +03 2026

ADIM 2 – Saldiri: PATH'in basina sahte/ klasoru ekleniyor
$ PATH="$PWD/sahte:$PATH" ./bin/linux/rapor
=== Aylik Satis Raporu ===
Rapor tarihi:
    !!! SAHTE 'date' calisti !!!
    Program PATH üzerinden kandirildi. ...

ADIM 3 – Duzeltilmis surum ayni saldiri altinda
=== Aylik Satis Raporu ===
Rapor tarihi: Sat Sep 19 19:11:32 +03 2026
```

The fake `date` here only prints a message. A real attacker, at the very same point, could do **anything** they wanted with **all the privileges** of whoever runs the program. If the program runs with elevated privilege (e.g. a system service), this is a direct **privilege escalation**.

The fixed code does three things at once: **1)** an absolute path, **2)** no shell, **3)** a small, known environment.

rapor_guvenli.c

```
char *const arguman[] = { "date", NULL };
char *const temiz_ortam[] = { "PATH=/usr/bin:/bin", "LANG=C", NULL };

pid_t pid;
int hata = posix_spawn(&pid, "/bin/date", NULL, NULL,
                      arguman, temiz_ortam);
```

rapor_guvenli.c

```
/* 1) C:\Windows\System32\hostname.exe - tam yol */
GetSystemDirectoryW(sistem, MAX_PATH);
swprintf(program, MAX_PATH, L"%ls\\hostname.exe", sistem);

/* 3) ortam bloğunda yalnız SystemRoot */
swprintf(ortam, MAX_PATH + 16, L"SystemRoot=%ls", kok);

/* 2) cmd.exe yok: program doğrudan başlatılır */
CreateProcessW(program, komut_satiri, NULL, NULL, FALSE,
               CREATE_UNICODE_ENVIRONMENT, ortam, NULL, &si, &pi);
```

✓ Rule

If you need to run another program: use an **absolute path**, do **not** run it through a shell (use `posix_spawn` / `execve` on Linux, `CreateProcessW` on Windows, instead of `system` / `popen`), and give the child process a **cleaned, known** environment. Best of all: avoid needing to run another program at all — get the date yourself with `time()` and `strftime()`, and the computer name with `gethostname()` or `GetComputerNameW()`.

11. Secure start-up: what to do in a program's first lines

Demo 1 showed how a single inherited environment variable (`PATH`) can turn all of a program's security upside down. `PATH` is not the only danger. The first chapter of our source book (Recipes 1.1–1.9) lists what a program should do in the **first lines** of its `main()` function. In this section we will update that list for today's systems and work through it step by step.

Program başlarken ebeveyninden ne devralır?

`main()`'in ilk satırlarında yapılacak iş budur

Ortam değişkenleri

`PATH`, `LD_PRELOAD`, `LD_LIBRARY_PATH` — hangi kod yüklenecek

Açık dosya tanıtıcıları

0/1/2 kapalıysa açtığınız dosya stdin olur

umask ve çalışma klasörü

oluşturduğunuz dosyanın izinleri buradan gelir

Sinyal düzeni ve kaynak sınırları

yok sayılan sinyal, düşük core limiti

Gerçek ve etkin kullanıcı kimliği

`setuid` ise ayrıcalık elinizdedir

Hiçbirine güvenmeyin: gerekeni kendiniz kurun, gerekmeyeni kapatın.

What does a program have when it starts?

When a process is started with `exec` (`CreateProcess` on Windows), it inherits the following from its **parent**:

Inherited	Why is it dangerous?	Recipe in the book
Environment variables	<code>PATH</code> , <code>LD_PRELOAD</code> , <code>LD_LIBRARY_PATH</code> change which program and library gets loaded	1.1
Open file descriptors	A secret file or socket the parent left open leaks into the child process; if 0–2 are closed, the first file opened becomes "standard output"	1.5
<code>umask</code>	Determines the permissions of new files; if <code>000</code> , files become writable by everyone	2.7
Working directory	Relative paths (<code>./ayar.ini</code>) can point into a folder the attacker chose	1.1
Resource limits	If core dumps are enabled, memory is written to disk the moment the process crashes	1.9
Identity and privileges	In a setuid program the real and effective user differ; more privilege than needed is carried along	1.3
Signal settings	Signals the parent ignores are also ignored in the child	13.5

The rule is simple: **do not trust anything the program itself did not set up**. The program's first job is to bring this inherited state to a **known, safe** state.

Step 1 – Clean the environment (Recipe 1.1)

There are two approaches: deleting the bad variables one by one (a **blocklist**), or wiping the environment entirely and rebuilding only what's needed (an **allowlist**). A blocklist is always incomplete, because you cannot know today which variable will turn out to be dangerous tomorrow. The right approach is the **allowlist**:

Linux – rebuild the environment with an allowlist

```
#include <stdlib.h>
#include <string.h>

static void ortami_temizle(void)
{
    const char *tz = getenv("TZ");          /* korunacak değişkeni önce kopyala */
    char tz_kopya[64] = "";
    if (tz && strlen(tz) < sizeof tz_kopya)
        strcpy(tz_kopya, tz);

    clearenv();                             /* glibc; taşınabilir değil (aşağıya bakın) */
    setenv("PATH", "/usr/bin:/bin", 1);     /* bilinen, yazılamaz dizinler */
    if (tz_kopya[0])
        setenv("TZ", tz_kopya, 1);
}
```

Why do we copy it first?

`getenv()` returns a pointer **into** the environment block. `clearenv()` invalidates that block; using the pointer afterward is a use of freed memory. Copy the value you want to keep into your **own buffer** first. `clearenv()` is glibc-specific; portable code uses `extern char **environ; environ = NULL;` instead.

If you **need** to use a variable (e.g. the user's `HOME`), treat it like an input: limit its length, verify it has the expected shape (Recipe 3.6). On Windows the environment is inherited the same way; there, alongside `PATH`, the **DLL search order** is also a danger (Step 6 below).

Step 2 — Tidy up file descriptors (Recipe 1.5)

There are two separate problems:

1. **If 0, 1, 2 are closed:** The attacker starts the program with its standard output closed. The first file the program opens (e.g. a password file) gets descriptor 1. Every subsequent `printf` then writes **into** that file.
2. **Extra open descriptors:** The parent hands the child a file or socket it left open without marking close-on-exec. The child can use that descriptor to reach a resource it normally couldn't.

Linux — guarantee 0–2, close the rest

```
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>

static void tanitici_duzenle(void)
{
    for (int fd = 0; fd <= 2; fd++) {
        if (fcntl(fd, F_GETFD) == -1) { /* kapalı mı? */
            int yeni = open("/dev/null", O_RDWR);
            if (yeni != fd) /* tam olarak o numarayı almalı */
                abort();
        }
    }
    /* Linux 5.9+ / glibc 2.34+: tek çağrıda 3 ve üstünü kapat */
    close_range(3, ~0U, 0);
}
```

On older systems, `close()` is called in a loop up to `sysconf(_SC_OPEN_MAX)`; if the limit is very large (e.g. a million) this can be slow, which is why `close_range` was added. For files you open yourself, always use the `O_CLOEXEC` flag, so they don't leak into a child process you launch later.

Step 3 — Restrict the permissions of new files (Recipe 2.7)

```
#include <sys/stat.h>
umask(077); /* yeni dosyalar yalnız sahibine açık: rw----- */
```

`umask` is a **mask**: the bits set in it are **subtracted** from the permissions you pass to `open()`. With `077`, no bits are left for group and others. Even so, write the permission **explicitly** when opening a sensitive file: `open(path, O_WRONLY | O_CREAT | O_EXCL, 0600)`. `O_EXCL` refuses to open the file if it already exists (e.g. if it's a link the attacker pre-created); recall the TOCTOU demo from [week 2](#).

Step 4 — Disable crash dumps (Recipe 1.9)

When a program crashes, the operating system may write the process's **entire memory** to a file to help with debugging. If that memory holds a password, a key, or decrypted data, all of it ends up sitting on disk as plain text. Demo 2's password can be found in exactly this way.

```
#include <sys/resource.h>
#include <sys/prctl.h>

static void dokumu_kapat(void)
{
    struct rlimit r = { 0, 0 };
    setrlimit(RLIMIT_CORE, &r);          /* çekirdek dökümü boyutu 0 */
    prctl(PR_SET_DUMPABLE, 0, 0, 0, 0); /* döküm yok, aynı kullanıcı bile ptrace ile bağlanamaz */
}
```

`PR_SET_DUMPABLE` provides a second benefit: another process running as the same user can no longer attach to this process as a debugger, and cannot read its `/proc/<pid>/mem`. We'll return to this line in [week 6](#) when we talk about debugger detection.

```
#include <windows.h>

static void dokumu_kapat(void)
{
    /* Çökme iletişim kutusunu ve kritik hata kutularını kapat */
    SetErrorMode(SEM_FAILCRITICALERRORS | SEM_NOGPFAULTERRORBOX);
}
```

On Windows, what actually determines whether a full memory dump is taken are Windows Error Reporting's (WER) local dump settings, which are configured at the system or user level. The most important thing a program can do from inside itself is: keep secrets in memory for the **shortest time possible** and erase them once the work is done (see sections 12 and 17). Even if a dump is taken, there should be no secret left inside it to find.

Step 5 — Drop unneeded privilege (Recipes 1.2–1.4)

If your program needs elevated privilege temporarily to do one thing (e.g. opening a port below 1024), it must **drop that privilege permanently** right after doing that thing.

Linux — permanently drop privilege in a setuid program, and verify it

```
#include <unistd.h>
#include <stdlib.h>

static void yetkiyi_birak(void)
{
    uid_t u = getuid();
    gid_t g = getgid();

    if (setresgid(g, g, g) != 0) abort(); /* önce grup: sonra buna yetkimiz kalmaz */
    if (setresuid(u, u, u) != 0) abort(); /* gerçek, etkin ve saklı kimliğin üçü de */

    uid_t r, e, s;
    getresuid(&r, &e, &s);
    if (r != u || e != u || s != u) abort(); /* bıraktığını DOĞRULA */
}
```

Three details are critical: (1) **group first, then user** – if you drop the user first, you no longer have the privilege to change the group. (2) `seteuid()`, which changes only the effective identity, is **temporary**; the **saved** identity stays privileged and privilege can be regained (week 2, Demo 13). (3) The result of dropping privilege is always **verified**; a `setuid()` call whose return value goes unchecked has caused real-world bugs.

On Windows the same idea is implemented with a **restricted token**: `CreateRestrictedToken` builds a token with administrator groups and privileges stripped out, and the risky work is done in a child process running with that token (Recipe 1.2). A stronger approach is **privilege separation** (Recipe 1.4): splitting the program into **two processes** – a small piece that needs privilege, and a large piece that runs unprivileged. A browser running every tab in an unprivileged "sandbox" process is today's version of this idea.

Step 6 – Care when loading other programs and libraries (Recipes 1.6–1.8)

Wrong	Why?	Correct
<code>system("convert " + dosya)</code>	The shell runs; a <code>;</code> , <code>&&</code> , <code>\$(...)</code> in the file name becomes a command (week 5)	<code>execve("/usr/bin/convert", argv, temiz_ortam)</code>
<code>execvp("convert", ...)</code>	Searches <code>PATH</code> (Demo 1)	<code>execve</code> with an absolute path
<code>CreateProcess(NULL, "C:\Program Files\Araç\A.exe", ...)</code>	An unquoted path with spaces: <code>C:\Program.exe</code> is tried first	Give <code>lpApplicationName</code> the full path, quote the command line
<code>LoadLibrary("yardimci.dll")</code>	A fake DLL placed in the app's or the working directory can be loaded during the search	<code>SetDefaultDllDirectories(LOAD_LIBRARY_SEARCH_DEFAULT_DIRS)</code> and an absolute path
<code>dlopen("libyardimci.so")</code>	<code>LD_LIBRARY_PATH</code> and search paths	An absolute path; <code>LD_*</code> variables should already be cleaned in Step 1

The Windows DLL search order problem is the **twin** of the `PATH` problem: if the program finds the library it wants in a folder the attacker left behind, that library's code runs with the program's privilege (CWE-427).

All together: `guvenli_baslat()`

guvenli_baslat.c – Linux skeleton (Windows equivalents in comments)

```
int main(int argc, char **argv)
{
    tanitici_duzenle(); /* Adım 2 – Windows: gerek yok, tanıtıcılar varsayılan olarak devralınmaz */
    ortami_temizle(); /* Adım 1 – Windows: SetDefaultDllDirectories + PATH doğrulaması */
    umask(077); /* Adım 3 – Windows: dosyayı açıkça kısıtlı DACL ile oluştur */
    dokumu_kapat(); /* Adım 4 – Windows: SetErrorMode */
    yetkiyi_birak(); /* Adım 5 – Windows: kısıtlanmış belirteç / ayrı süreç */

    /* ... programın asıl işi ... */
    return 0;
}
```

The order is not arbitrary: descriptors and the environment must be tidied up **before the program does anything else**, because every later step (e.g. writing an error message) depends on them.

🔥 How does an evaluator test this?

They start the program with standard output closed, with a poisoned `PATH` and `LD_PRELOAD`, with `umask 000`; they trigger a crash and look for a dump file; for `setuid` programs they check whether privilege was really dropped by reading the `Uid:` line in `/proc/<pid>/status`. On Windows they drop a fake DLL into the app's folder and the working directory and see whether it gets loaded.

🔧 How is this applied in the field?


In mobile payment libraries, the native layer checks whether the `LD_PRELOAD` variable is set; it even writes a random variable with `setenv` and reads it back with `getenv` to detect whether the `getenv` function itself has been hooked. If the value doesn't match, someone is intercepting and altering the call. This is the point where secure start-up turns into "runtime self-protection" (week 6).

12. Secrets left in memory

What happens to a password after you use it? The variable goes out of scope, but the **memory is not erased**; those bytes sit there until something else is written over them. If an attacker can see that memory, they read the secret. There are more situations where memory can be seen than you might think:

Sır bellekten nasıl dışarı sızır?

değişken kapsam dışına çıkar, baytlar kalır



Çözüm: `explicit_bzero` / `SecureZeroMemory` — derleyicinin kaldıramayacağı silme.

- The **crash dump** (core dump) file created when the program crashes,
- The **swap area** and the **hibernation file**, where the operating system writes memory to disk,
- A **debugger** attaching to the process,
- Another piece of code that later reuses the same memory reading its **old contents**.

Demo 2 – The password left in memory

Demo 2 · `code/week-01/02-bellekte-parola` · CWE-14 (compiler removal of code to clear buffers), CWE-316

The program reads the password from `parola.txt`, computes a key with it, then tries to erase the password. The demo dumps the program's memory to a file and searches it for the password: on Linux with `gdb`'s `gcore` command, on Windows by having the program write its own memory with `MiniDumpWriteDump`. There are three versions: one that **never erases**, one that erases with `memset`, and one that erases with `explicit_bzero` / `SecureZeroMemory`.

parola.c (excerpt)

```
unsigned long giris_yap(const char *dosya)
{
    char parola[64];
    /* düşük düzey okuma: parola yalnız bu dizide durur */
    int fd = dosya_ac(dosya);
    int n = (int)dosya_oku(fd, parola, sizeof(parola) - 1);
    ...
    unsigned long anahtar = anahtar_turet(parola, (size_t)n);

    #if SILME == 1
        memset(parola, 0, sizeof(parola));          /* "ölü yazma" */
    #elif SILME == 2
        #ifdef _WIN32
            SecureZeroMemory(parola, sizeof(parola)); /* kaldırılamaz */
        #else
            explicit_bzero(parola, sizeof(parola));   /* kaldırılamaz */
        #endif
    #endif
    return anahtar;
}
```

```
cd code\week-01\02-bellekte-parola
.\demo.ps1
```

```
cd code/week-01/02-bellekte-parola
sh demo.sh
```

sh demo.sh – output (the result is the same on Windows)

```
ADIM 1 – Parolayı hic silmeyen surum
SONUC: parola dokumde 1 yerde BULUNDU -> Gizli-Parola-2026

ADIM 2 – memset ile silen surum (-02)
SONUC: parola dokumde 1 yerde BULUNDU -> Gizli-Parola-2026

ADIM 3 – explicit_bzero ile silen surum (-02)
SONUC: parola dokumde bulunamadi

Neden? giris_yap() fonksiyonunun makine kodu:
memset surumunde silme komutu sayisi -> 0
explicit surumunde explicit_bzero cagrisi -> 1
```

The second step is surprising: the code plainly writes `memset`, yet the password is **still in memory**! The reason is the compiler's optimisation. The compiler reasons like this: *"the `parola` array is never read after this line; writing zeros into it*

has no effect on the program's result. So this write is redundant, I can remove it." This is called **dead store elimination**. This optimisation, which speeds up ordinary programs, leaves a secret sitting in memory in security code. This isn't unique to GCC either: MSVC with `/O2` on Windows produces the same result. See for yourself: on Linux, search for `giris_yap` with `objdump -d bin/linux/parola_memset | less`; in Visual Studio, put a breakpoint on `giris_yap` and open **Debug > Windows > Disassembly**. There isn't a single instruction belonging to `memset`.

✓ Rule

To erase secrets, use functions the compiler **cannot** remove: `explicit_bzero` on Linux/glibc, `SecureZeroMemory` on Windows, `memset_s` on compilers that support C11 Annex K, `OPENSSL_cleanse` in OpenSSL. Extra layers: disabling crash dumps (`setrlimit(RLIMIT_CORE, 0)`), preventing memory from being written to disk (`mlock`), keeping the secret for the shortest time possible, and not multiplying copies of it.

🔧 How is this applied in the field?


In software that needs high security, such as mobile payment, the one-time payment key is only decrypted at the moment of payment, and only in the native layer; right after use, the buffer is erased by overwriting it with random data. On the Java side, sensitive data is kept in erasable `byte[]` arrays instead of the unerasable `String`. An independent evaluator tests this by dumping memory during and after the payment and searching for the key — exactly what we did in Demo 2.

13. Process memory: where does your data live?

To understand overflow bugs, we first need to know how a program's memory is laid out. A process running on Linux has virtual memory laid out roughly like this:

Süreç belleği: verileriniz nerede duruyor?

yığın aşağı, öbek yukarı büyür — ortada boşluk



Taşmanın sonucu hangi bölgede olduğuna bağlıdır: yığında dönüş adresi, öbekte komşu blok.

When a function is called, a **stack frame** is allocated for it on the stack. Local variables sit **side by side** in this frame. C does not check when the end of an array is reached: if you write a 17th byte into an `ad[16]` array, C does not stop you, and that byte overwrites whatever comes right after it in memory.

Let's look at the structure from Demo 3:

```
struct oturum {
    char ad[16];
    int yönetici; /* 0 = normal kullanıcı */
};
```

Yapının bellekteki hâli: taşma neyi ezer?

`ad[16]` ve `yönetici` yan yana durur

ÖNCE

"ayse"
`ad[16]`

00 00 00 00
`yönetici = 0`

SONRA

"AAAAAAAAAAAAAAAA"
17 karakterlik "AAAA...AAB" kopyalandı

B 00 00 00
`yönetici ≠ 0`

Tek bayt taşma yetkiyi değiştirdi: sınır denetimi olmayan kopya, komşu alanı ezer.

x86-64 processors are **little-endian**: an integer's **least significant byte** sits at the **lowest** address. That's why writing `'B'` (0x42) into the first byte of `yönetici` turns the integer into 66 — and any non-zero value means "administrator."

14. Buffer overflows: how they happen, how to prevent them

In the previous section we saw the regions of a process's memory. Now let's meet the oldest, and still the most common, bug in that memory. A **buffer overflow** is when a program writes **more data into an array than the space allocated for it**. The overflowing bytes don't vanish; they get written into the memory right next to the array — that is, on top of **other variables**. "Out-of-bounds write" (CWE-787) has sat near the top of the CWE Top 25 list for years.

Why does C overflow so easily?

In languages like Java, C#, and Python, when you try to write past the end of an array, the runtime **stops you** (`ArrayIndexOutOfBoundsException`). In C and C++, however:

1. An array is nothing more than consecutive bytes in memory; **it does not know its own size**.
2. When you pass an array to a function, only its **starting address** goes along (`char *`); the size information is lost.
3. Strings are not a separate type; they are a `char` array marked at the end with a `'\0'` byte. Functions like `strcpy`, `strcat`, and `sprintf` copy the source **until they see the zero byte**; they never ask how much room the destination has.
4. The compiler adds **no bounds checking**, so as not to give up speed.

The result: checking the bound is **entirely the programmer's responsibility**. Forget it once, and the bug is there.

A stack overflow, step by step

The function below copies a username into a small local array:

Buggy: the source length is never checked

```
void selamla(const char *ad)
{
    int yetkili = 0;
    char tampon[16];

    strcpy(tampon, ad);          /* ad 16 bayttan uzunsa? */
    printf("Merhaba %s\n", tampon);
    if (yetkili) { /* ... */ }
}
```

When the function is called, a **stack frame** is created for it. Inside that frame sit the local variables, the address of the previous frame, and the **return address** that says where to go once the function ends. A simplified view (lowest address at the top; the real layout varies by compiler):

Yığında taşma adım adım

strcpy aşağıdan yukarı yazar



Kanarya, dönüş adresinden önce durur: ezilirse program sonlanır (ama yetkili zaten ezilmiştir).

There are three cases:

Input length	What happens?	Result
≤ 15 characters	Everything fits in the buffer	The program works correctly
A little more	The overflowing bytes write over <code>yetkili</code> ; <code>yetkili</code> is no longer 0	The logic breaks: an unauthorized user is treated as authorized (Demo 3)
A lot more	The return address is corrupted; once the function ends, the processor jumps to a meaningless address	The program crashes (denial of service); on an unprotected system, control flow can be hijacked

The last part of the third row is what makes an overflow this dangerous: whoever controls the return address controls **which code the program runs next**. This was the entry point for most outbreaks from the Morris worm in 1988 to Slammer in 2003 ([week 2](#)). Today's operating systems and compilers add protections that make this path much harder (below), but the bug itself is **still in your code**, and the protections don't close every case. In this course we care about how the bug **happens, is found, and is prevented** — not about how to exploit it.

Types of overflow

Type	Where?	Typical cause	CWE
Stack overflow	A local array	<code>strcpy</code> , <code>gets</code> , unbounded <code>scanf("%s")</code>	CWE-121
Heap overflow	A block allocated with <code>malloc</code>	A <code>memcpy</code> with a miscalculated length; heap management metadata is corrupted	CWE-122
Off-by-one	Any array	<code><=</code> instead of <code><</code> ; forgetting to reserve room for the <code>'\0'</code> terminator	CWE-193
Out-of-bounds read	Any array	Trusting a length the other party reports	CWE-125
Integer-driven	Any array	<code>n * size</code> overflows, a small block is allocated, then a large copy is made into it	CWE-190 → 787



An out-of-bounds read is as dangerous as a write: Heartbleed (2014)

In OpenSSL's "heartbeat" extension, the client sent a message and that message's **length**; the server sent the message back as-is. The server never checked whether the declared length was really the length of the message it received. When the client sent a 1-byte message but said "64 KB," the server sent back 64 KB from its own memory — which could contain other users' passwords, session data, even the server's **private key** (CVE-2014-0160). Nothing was written, nothing crashed; it was **only read**. The fix was a single check: "if the declared length is bigger than the received record's length, drop the message."

Off-by-one: the most innocent-looking overflow

Buggy: no room for the terminator

```
char ad[8];
for (int i = 0; i <= 8; i++)    /* 9 kez döner: ad[8] dizinin dışında */
    ad[i] = kaynak[i];
```

Buggy: strncpy does not guarantee a terminator

```
char ad[8];
strncpy(ad, kaynak, sizeof ad); /* kaynak ≥ 8 ise ad '\0' ile BİTMEZ */
printf("%s\n", ad);             /* printf dizinin sonundan okumaya devam eder */
```

`strncpy`'s name is misleading: it is **not** a safe `strcpy`. If the source is longer than the destination it does not add a terminator; if it's shorter, it needlessly pads the rest of the destination with zeros. Even a **single-byte** overflow is enough

to change the low byte of an adjacent variable, or a saved frame address.

Dangerous functions and what to use instead

Dangerous	Problem	Instead	Note
<code>gets(s)</code>	Never knows the bound	<code>fgets(s, sizeof s, stdin)</code>	Removed from the standard in C11
<code>strcpy(d, s)</code>	No target size	Check the length first, then <code>memcpy</code> ; or <code>snprintf(d, sizeof d, "%s", s)</code>	<code>strncpy</code> on glibc 2.38+ and the BSDs
<code>strcat(d, s)</code>	No target size	Compute the remaining room; concatenate with <code>snprintf</code>	<code>strlcat</code>
<code>sprintf(d, ...)</code>	Output length unknown	<code>snprintf(d, sizeof d, ...)</code> and check the return value	If the return is $\geq$ size, it means truncation
<code>scanf("%s", s)</code>	Unbounded read	<code>scanf("%15s", s)</code> (size - 1) or <code>fgets</code>	
<code>strncpy(d, s, n)</code>	No <code>'\0'</code> guarantee	<code>snprintf</code> or <code>strncpy</code>	
<code>memcpy(d, s, n)</code>	If <code>n</code> is not validated	Check <code>n ≤ target size</code> first	Most common bug: <code>n</code> comes from outside
<code>alloca(n)</code> , VLA <code>char a[n]</code>	A large <code>n</code> overruns the stack	A fixed upper bound, or <code>malloc</code>	

The modern equivalent of Recipes 3.3 and 3.4

The book recommends handing string operations off to a bounds-checked library (SafeStr) to prevent overflows. The idea still holds today, the tools have changed: `snprintf` and `strncpy` in C, the `strcpy_s` / `strcat_s` family from C11 Annex K on the Microsoft compiler (glibc does not support this family), and `std::string`, `std::vector`, `std::array::at()`, and `std::span` in C++. Microsoft's secure development process uses a "banned functions" header file that marks dangerous functions as a **build error**; you can use the same method in your project.

The right pattern: validate first, then copy with a bound

Fixed: length check + bounded copy + always a terminator

```
#include <stdio.h>
#include <string.h>

int selamla(const char *ad)
{
    char tampon[16];
    size_t n = strlen(ad, sizeof tampon);

    if (n == sizeof tampon)           /* 15 karakterden uzun: reddet, sessizce kesme */
        return -1;

    memcpy(tampon, ad, n);
    tampon[n] = '\0';
    printf("Merhaba %s\n", tampon);
    return 0;
}
```

This pattern makes three decisions:

1. With `strlen`, the length is measured over **at most** the buffer size in bytes; even if the source never terminates, the read never goes past the bound.
2. Long input is **rejected**. Silently truncating is sometimes correct (e.g. a title to be displayed on screen), but when a credential, a file path, or a command is truncated, its meaning can change: `/home/ayse/gizli_rapor.txt` truncated can become `/home/ayse/gizli`. From a security standpoint, **rejecting** ambiguous input is safer.
3. The terminator is written **explicitly**, in every case.

In C++, the language does the same job for you:

C++ – the type carries its own size

```
#include <string>
#include <iostream>

void selamla(const std::string &ad)
{
    if (ad.size() > 15) throw std::invalid_argument("ad cok uzun");
    std::string tampon = ad;           /* kendi belleğini kendi yönetir */
    std::cout << "Merhaba " << tampon << '\n';
}
```

⚠ C++ doesn't always protect you either

`std::vector::operator[]` does not check bounds; `at()` does. If you take the pointer from `std::string::c_str()`, hand it to a C function, and write to it there, all of C's problems come right back. Using a safe type does not exempt you from thinking about bounds.

Who catches the overflow? Layers of defence

Defence against overflows is layered too. The **innermost** layer is correct code; the **outer** layers **limit the damage** when a bug slips through:

Layer	Tool	What it does	When?
Code	Bounds checking, safe API	Never lets the overflow happen at all	Always
Compiler warning	<code>-Wall -Wextra -Wformat-security, MSVC /W4 /sd1</code>	Flags suspicious calls at compile time	Development
Static analysis	<code>clang --analyze, cppcheck, MSVC /analyze</code>	Looks for possible overflows without running the code	Development, CI
Sanitizer	AddressSanitizer (<code>-fsanitize=address, MSVC /fsanitize=address</code>)	Checks every memory access; stops the program at the first overflow and says where	Testing
Fuzzing	libFuzzer, AFL++	Feeds the program millions of random inputs	Testing (week 4)
Compiler protection	Stack canary (<code>-fstack-protector-strong, MSVC /GS, _FORTIFY_SOURCE</code>)	Places a "canary" in front of the return address; terminates the program if it's corrupted	Release
Operating system	DEP/NX, ASLR	Blocks running code in a data region; changes addresses on every run	Release
Hardware	Shadow stack (Intel CET, ARM PAC)	Keeps a separate, protected copy of the return address	Release (newer processors)

We will examine each of these protections [in week 4](#) by turning them on and off one at a time. This week keep this idea in mind: **a protection makes exploitation harder; it does not fix the bug**. When the canary is corrupted, it terminates the program — meaning the attacker can still crash the program even if they can't hijack control flow. That, too, is a denial of service.

Catching an overflow with AddressSanitizer

If you compile the buggy `selamla` function above with the sanitizer turned on and call it with a 20-character name, the overflow is caught before the program's logic is even corrupted. The output looks like this (abridged; addresses change on every run). You will produce a report of the same kind yourself with Demo 3's `giris_asan` target:

```
==12345==ERROR: AddressSanitizer: stack-buffer-overflow on address 0x7ffd...
WRITE of size 21 at 0x7ffd... thread T0
    #0 in strcpy
    #1 in selamla ornek.c:7
    #2 in main ornek.c:15
Address 0x7ffd... is located in stack of thread T0 at offset 48 in frame
    #0 in selamla ornek.c:3
This frame has 1 object(s):
    [32, 48) 'tampon' <== Memory access at offset 48 overflows this variable
```

Read it in this order: **what** happened (a stack overflow, a 21-byte write), **where** it happened (`ornek.c` line 7, `strcpy`), **which variable** (`tampon` , 16 bytes: the 32–48 range; the access is at 48, i.e. it starts exactly one byte past the bound). The sanitizer slows the program down roughly twofold and increases memory use; that's why it is used in **test** builds, not in release builds.

How does an evaluator test this?

In source code review they first search for dangerous functions (`grep -nE "strcpy|strcat|sprintf|gets|scanf"`), then check whether every externally supplied length field is validated before use. In dynamic testing they feed every input an overly long value, a zero length, a negative length, and data shorter than what the field declares; they run the program under a sanitizer and do fuzzing. On the binary they check whether the stack canary, NX, and ASLR are turned on (`checksec` on Linux, `dumpbin /headers` on Windows).

How is this applied in the field?

In security-sensitive native libraries, the rules of defensive programming are written down at the very start of the guide: every input and output is validated, no data crosses an interface in the clear, assets are kept in memory for the shortest time possible, and memory is wiped with a secure erase function. Some teams use their own versions of standard functions like `memcpy`, `memset`, `strcmp` instead of the originals. There are two reasons for this: to make it harder for an attacker to hook these well-known functions and read data in transit, and to bring the code of these functions itself inside the scope of integrity checking (week 6).

15. Demo 3 – Privilege escalation via buffer overflow

Demo 3 · `code/week-01/03-tasma-giris` · CWE-121 (stack-based overflow), CWE-787 (out-of-bounds write)

giris.c (buggy)

```
struct oturum o;
o.yonetici = 0;

strcpy(o.ad, argv[1]); /* HATA: hedefin 16 bayt olduğu hiç denetlenmiyor */

if (o.yonetici)
    printf(">>> YONETICI paneline erisim verildi! <<<\n");
```

The demo compiles the same buggy code in four different ways and tries the same attack:

Version	Linux / WSL (GCC)	Windows (MSVC)	What do we expect to see?
<code>giris</code>	<code>-00</code> , no protection	<code>/Od</code>	The attack succeeds
<code>giris_asan</code>	<code>-fsanitize=address</code>	<code>/fsanitize=address</code>	Does AddressSanitizer catch it?
<code>giris_denetimli</code>	<code>-02 -D_FORTIFY_SOURCE=2</code>	<code>/02 + strcpy_s</code>	Does the library's size check catch it?
<code>giris_guvenli</code>	Fixed code	Fixed code	The input is rejected

```
cd code\week-01\03-tasma-giris
.\demo.ps1
```

```
cd code/week-01/03-tasma-giris
sh demo.sh
```

sh demo.sh – output

```
ADIM 1 – Normal giris
$ ./bin/linux/giris ayse
Hos geldin, ayse
yonetici alani = 0 (0x00000000)
Normal kullanıcı oturumu.

ADIM 2 – Saldiri: 16 bayttan 1 bayt fazla ad (17 karakter)
$ ./bin/linux/giris AAAAAAAAAAAAAAAB
Hos geldin, AAAAAAAAAAAAAAAB
yonetici alani = 66 (0x00000042)
>>> YONETICI paneline erisim verildi! <<<

ADIM 3 – Ayni saldiri, AddressSanitizer ile derlenmis surumde
$ ./bin/linux/giris_asan AAAAAAAAAAAAAAAB
Hos geldin, AAAAAAAAAAAAAAAB
yonetici alani = 66 (0x00000042)
>>> YONETICI paneline erisim verildi! <<<
^ ASan sessiz kaldı: tasma yapinin ICINDE kaldı.

ADIM 4 – Yapinin disina tasan uzun ad, ASan surumu
==1750==ERROR: AddressSanitizer: stack-buffer-overflow on address ...
WRITE of size 41 at 0x7fffffffe8c4 thread T0

ADIM 5 – Ayni kisa saldiri, _FORTIFY_SOURCE=2 ile derlenmis surumde
$ ./bin/linux/giris_denetimli AAAAAAAAAAAAAAAB
*** buffer overflow detected ***: terminated
(program, derleyicinin ekledigi denetimle durduruldu)

ADIM 6 – Duzeltilmis surum
$ ./bin/linux/giris_guvenli AAAAAAAAAAAAAAAB
Reddedildi: ad 1-15 karakter; yalniz harf, rakam, _ ve - olabilir.
```

The flow is the same on Windows. Only in step 5, `strcpy_s` knows the target is 16 bytes, so it sees the overflow and stops the program:

.\demo.ps1 – step 5

```
> .\bin\windows\giris_denetimli.exe AAAAAAAAAAAAAAAB
*** Guvenli kutuphane islevi tasmayi algiladi: program durduruldu ***
(cikis kodu: 3)
```

Three lessons come out of this demo:

1. **A single byte of overflow is enough.** The attacker doesn't need to crash the program; changing a single flag can be enough to escalate privilege.
2. **Tools have limits.** AddressSanitizer places a "forbidden zone" around every object in memory; but because the overflow stayed **inside** the same struct (`ad` → `yonetici`), it didn't catch it. It caught it immediately once the overflow went **outside** the struct. `_FORTIFY_SOURCE` (Linux) and `strcpy_s` (Windows), on the other hand, knew the target (`o.ad`) was 16 bytes, so they caught the short overflow too. **Don't rely on a single tool.**
3. **The real fix lives in the code.** Tools find the bug; fixing it is the programmer's job.

giris_guvenli.c (excerpt)

```
/* İzin listesi: 1-15 karakter; harf, rakam, '_' ve '-' */
static int ad_gecerli_mi(const char *s)
{
    size_t n = strlen(s, 16);
    if (n == 0 || n >= 16)
        return 0;
    for (size_t i = 0; i < n; i++) {
        unsigned char c = (unsigned char)s[i];
        if (!isalnum(c) && c != '_' && c != '-')
            return 0;
    }
    return 1;
}
...
if (!ad_gecerli_mi(argv[1])) { /* reddet */ }
struct oturum o = { .yonetici = 0 }; /* güvenli varsayılan */
snprintf(o.ad, sizeof(o.ad), "%s", argv[1]); /* boyutu bilen kopya */
```



What happens in real attacks?


Here, the overflowing bytes changed a flag. The function's **return address** also sits on the stack, right after the local variables; if the overflow reaches that far, the attacker can redirect the program's control flow. Modern systems take layered countermeasures against this: the **stack canary** (a secret value placed before the return address that gets corrupted by an overflow), **DEP/NX** (forbids executing data on the stack as code), and **ASLR** (changes addresses on every run). We'll cover these in detail in [week 4](#); what matters here is knowing that they **don't fix the bug**, they only **make exploitation harder**.

16. Integers and the signed-length bug

In C, `int` is **signed**, while `size_t` is **unsigned**. Computers store negative numbers in **two's complement**: the 64-bit representation of `-1` is `0xFFFFFFFFFFFFFFFF`. Read as unsigned, that same bit pattern is **18,446,744,073,709,551,615**. Functions like `memcpy`, `malloc`, and `read` take a `size_t` for their size parameter; so the moment a negative `int` is passed to them, it turns into a **gigantic** number.

İşaretsiz uzunluk hatası: aynı bitler, iki farklı sayı

int işaretsiz, size_t işaretsiz — dönüşüm sessizdir



Kural: boyutlar baştan sona size_t tutulur; -Wsign-conversion açılır.

Demo 4 — The signed-length bug

Demo 4 · `code/week-01/04-isaretli-uzunluk` · **CWE-195 (signed-to-unsigned conversion error)**

kopya.c (buggy)

```
#define KAYIT_BOYUTU 16

static int kaydi_kopyala(char *hedef, const char *kaynak, int uzunluk)
{
    /* HATA: alt sınır (negatif değer) denetlenmiyor */
    if (uzunluk > KAYIT_BOYUTU)
        return -1;
    memcpy(hedef, kaynak, uzunluk); /* int -> size_t: -1 -> 2^64 - 1 */
    return 0;
}
```

```
cd code\week-01\04-isaretli-uzunluk
.\demo.ps1
```

```
cd code/week-01/04-isaretli-uzunluk
sh demo.sh
```

sh demo.sh — output

```
ADIM 0 – Derleyici bu hatayı görüyor mu?
$ gcc -Wall -Wextra -fsyntax-only -I.././common kopya.c
-> Uyarı yok: -Wall -Wextra bu donusumu bildirmez.
$ gcc -Wsign-conversion -fsyntax-only -I.././common kopya.c
kopya.c:21:27: warning: conversion to 'size_t' from 'int'
      may change the sign of the result [-Wsign-conversion]
```

```
ADIM 1 – Normal kullanım
$ ./bin/linux/kopya 8
Kopyalandı: ORNEK-KA
```

```
ADIM 2 – Çok büyük uzunluk: denetim yakalar
$ ./bin/linux/kopya 100
Reddedildi: uzunluk çok büyük.
```

```
ADIM 3 – Saldırı: negatif uzunluk denetimi atlatır
$ ./bin/linux/kopya -1
Segmentation fault
(cikis kodu: 139 - 139 = bellek hatasi, SIGSEGV)
```

```
ADIM 4 – Aynı saldırı, AddressSanitizer ile
==1823==ERROR: AddressSanitizer: negative-size-param: (size=-1)
```

```
ADIM 5 – Düzeltilmiş sürüm
$ ./bin/linux/kopya_guvenli -1 ; ./bin/linux/kopya_guvenli 12abc ; ./bin/linux/kopya_guvenli 8
Reddedildi: uzunluk 0-16 arasında bir tamsayı olmalı.
Reddedildi: uzunluk 0-16 arasında bir tamsayı olmalı.
Kopyalandı (8 bayt): ORNEK-KA
```

On Windows you see two differences. In step 3 the program ends with code `0xC0000409`: Windows notices the stack is corrupted and terminates the process immediately. In step 0, MSVC's C compiler gives **no warning at all**, even with `/Wall`; the same code compiled as C++ does produce a warning:

`.\demo.ps1 – step 0`

```
> cl /Zs /W4 kopya.c (C olarak derle)
-> Uyarı yok: MSVC'nin C derleyicisi bu donusumu /Wall ile bile bildirmez.
> cl /Zs /W4 /w44365 /TP kopya.c (ayni kod, C++ olarak)
kopya.c(21): warning C4365: 'argument': conversion from 'int' to 'size_t', signed/unsigned mismatch
```

The check asks "is `uzunluk > 16`?" ; `-1` answers "no" to that question, so it gets through. A length field in a packet coming over the network, a record size in a file header — **any number the attacker controls** can lead to this bug. The compiler **could** have warned us, but only if we asked: `-Wsign-conversion` is not in GCC's default warnings, and MSVC's C compiler doesn't have it at all. Knowing which warnings are on, and reading them, is a **free** security tool; a tool's silence does not mean the code is correct.

`kopya_guvenli.c (excerpt)`

```
static int uzunluk_oku(const char *metin, size_t *sonuc)
{
    char *son;
    errno = 0;
    long deger = strtol(metin, &son, 10);
    if (errno != 0 || son == metin || *son != '\0')
        return -1; /* sayı değil ya da taşıtı */
    if (deger < 0 || deger > KAYIT_BOYUTU)
        return -1; /* hem alt hem üst sınır */
    *sonuc = (size_t)deger;
    return 0;
}
```

✓ Rule

Keep sizes and lengths in `size_t` from the start. Validate every externally supplied number against **both a lower and an upper bound**. Use `strtol`, which reports errors, instead of `atoi`. Compile with `-Wall -Wextra -Wconversion -Wsign-conversion` turned on.

We will go deeper into this wider family of integer bugs — overflow, sign conversion, boundary values — [in week 4's undefined-behaviour section](#).

17. Memory management and security

A buffer overflow happens when we forget **how much** we're writing. The bugs in this section happen when we forget **who owns** a piece of memory and **how long** it stays valid. In C and C++, memory allocated from the **heap** lives until the programmer frees it; the language does not track this lifetime for you.


The four lifetimes of memory

Lifetime	How is it created?	When does it end?	Typical bug
Static	A global variable, a <code>static</code> local variable	When the program ends	Multiple threads writing to it at the same time
Automatic	A local variable inside a function	When the function returns	Returning the address of a local variable
Dynamic	<code>malloc</code> / <code>calloc</code> / <code>new</code>	When <code>free</code> / <code>delete</code> is called	Leak, double free, use-after-free
Thread-local	<code>_Thread_local</code> / <code>thread_local</code>	When the thread ends	—

The lifecycle of dynamic memory has four steps: **allocate** → **validate** → **use** → **free**. A distinct class of bug lives at each step.

Belleğin dört ömrü ve her adımın CWE'si

her adımın atlanması ayrı bir zafiyet sınıfı üretir



Altı adımın hepsi yapılmazsa bellek hatası kaçınılmazdır.

Classes of error

Error	What happens?	Security consequence	CWE
Memory leak	Allocated memory is never freed	A long-running server runs out of memory → denial of service; secrets stay in memory longer than they should	401
Double free	The same block is <code>free</code> d twice	The memory manager's internal lists get corrupted; the next two <code>malloc</code> calls can hand the same block to two different objects	415
Use-after-free (UAF)	A <code>free</code> d block is accessed through a stale pointer	If the block has been handed to another object, the stale pointer reads or modifies someone else's data	416
Uninitialized read	A value is read before it's assigned	A previous user's data (e.g. a password) leaks; behaviour becomes unpredictable	457, 908
No NULL check	<code>malloc</code> fails, the result is not checked	A crash; on some systems, a write to a low address	476
Mismatch	Something allocated with <code>new[]</code> is freed with <code>delete</code> , something allocated with <code>malloc</code> is freed with <code>delete</code>	Undefined behaviour; the memory manager gets corrupted	762
Size-calculation overflow	<code>n * sizeof(T)</code> overflows	A small block is allocated, a large write follows → heap overflow	190, 131
Secret not erased	A freed block is not zeroed	The block gets handed to other code by the next <code>malloc</code> ; the secret inside it can be read	226, 244

Use-after-free: why is it so dangerous?

Buggy: two pointers point at the same block

```
typedef struct { char ad[32]; int yetki; } Kullanici;

Kullanici *aktif = malloc(sizeof *aktif);
Kullanici *onbellek = aktif;          /* ikinci bir sahip daha */
/* ... */
free(aktif);                          /* oturum kapandı */
aktif = NULL;                         /* aktif güvende, ama onbellek hâlâ eski adresi tutuyor */

char *not = malloc(sizeof(Kullanici)); /* bellek yöneticisi aynı bloğu yeniden verebilir */
strcpy(not, "....");                  /* not yazılıyor */

if (onbellek->yetki) { /* ... */      /* UAF: artık 'not'un baytlarını okuyor */
```

`free` does not return memory to the operating system; it puts it on the memory manager's **"free" list**. The next `malloc` of the same size will most likely get the same block back. The stale pointer (`onbellek`) is now pointing at a **completely different object**. The program's logic reads bytes that were written for some other purpose as "the user's privilege." This is why UAF bugs are one of the largest classes of browser and operating system vulnerabilities.

The `aktif = NULL;` line in the example is a **good habit**, but it only protects **that one pointer**. The real problem is that **more than one pointer** claims ownership of the same block. The solution is an **ownership rule**:

✓ The ownership rule

Every dynamic block has **exactly one owner**, and only that owner frees it. Other code **borrow**s the block and never stores it anywhere that will outlive the owner. If ownership passes from one function to another, that is written explicitly in the function's name and documentation (e.g. `..._al()` (take) takes over ownership, `..._goster()` (show) only borrows it).

We will exploit a UAF step by step in C/C++ and catch it with AddressSanitizer [in week 4](#).

Fixed versions of the other errors

NULL check and overflow-free size calculation

```
/* Hatalı: n büyükse n * sizeof(Kayit) taşar, küçük bir blok ayrılır */
Kayit *k = malloc(n * sizeof(Kayit));

/* Doğru 1: calloc boyut çarpımında taşmayı denetler ve belleği sıfırlar */
Kayit *k = calloc(n, sizeof(Kayit));
if (k == NULL) return HATA_BELLEK;

/* Doğru 2: yeniden boyutlandırmada reallocarray (glibc 2.26+, BSD) */
Kayit *yeni = reallocarray(k, n2, sizeof(Kayit));
if (yeni == NULL) { free(k); return HATA_BELLEK; } /* k'yi kaybetme */
k = yeni;
```

The realloc trap

```
/* Hatalı: realloc başarısız olursa NULL döner ve eski blok hâlâ ayrılmış durumdadır;
   onu gösteren tek işaretçinin üzerine yazdık → sızıntı */
tampon = realloc(tampon, yeni_boyut);
```

`realloc` has a second hidden problem: if the block **moves**, the old block is freed but **not zeroed**. If it held a secret, the old copy sits somewhere in memory forever. Don't use `realloc` for buffers that hold secrets; allocate a new block, copy into it, and **erase** the old one before freeing it.

Erasing secrets from memory (Recipe 13.2)

Demo 2, in the "Secrets left in memory" section, showed that the compiler can see the line `memset(parola, 0, ...)` as "a write to somewhere that's never read afterward" (a dead store) and **remove it entirely**. To prevent this, functions the compiler will **not** remove are used:

Platform	Function	Note
Linux (glibc 2.25+), BSD	<code>explicit_bzero(p, n)</code>	The most common choice
Windows	<code>SecureZeroMemory(p, n)</code>	In <code>windows.h</code>
C11 Annex K	<code>memset_s(p, smax, 0, n)</code>	Not available on every compiler
C23	<code>memset_explicit(p, 0, n)</code>	New standard
If none is available	Writing byte by byte through a <code>volatile</code> pointer	The portable fallback

cen429_sil — portable secure erasure

```
#include <string.h>
#ifdef _WIN32
# include <windows.h>
#endif

static void cen429_sil(void *p, size_t n)
{
    #if defined(_WIN32)
        SecureZeroMemory(p, n);
    #elif defined(__GLIBC__) || defined(__OpenBSD__) || defined(__FreeBSD__)
        explicit_bzero(p, n);
    #else
        volatile unsigned char *v = (volatile unsigned char *)p;
        while (n--) *v++ = 0;
    #endif
}
```

Where the erasure happens matters too. A secret can have more copies than you think:

- Structs passed **by value** to a function (a copy is left on the stack on every call).
- The old block left behind when `realloc` moves it.
- `std::string`'s small-string optimisation (SSO) and its copies; `std::string` does **not** erase its contents when it's destroyed.
- I/O buffers (the `stdin` buffer that `fgets` reads into).

- The operating system's **swap file** and **hibernation file** (below).

How is this applied in the field?

In sensitive libraries, session data is erased by filling it with **random values** the moment the operation ends; on the native side, this filling is the **only** place the random number generator is used. A random value instead of zero also removes the tell-tale sign in memory that "a secret was just here" (a long run of zero bytes). The "creation → deletion" column in the asset table is written in enough detail to show exactly **which line** each asset is erased on.

Preventing memory from being written to disk (Recipe 13.3)

Under memory pressure, the operating system writes rarely used pages to the swap area (Linux swap, Windows `pagefile.sys`). When the computer enters hibernation, **all** of memory is written to disk. No matter how well you erase your secret from memory, a copy may already have been written to disk in the meantime. The countermeasure is to **lock** the pages holding secrets:

```
#include <sys/mman.h>

unsigned char anahtar[32];
if (mlock(anahtar, sizeof anahtar) != 0) {
    /* RLIMIT_MEMLOCK aşılmış olabilir; hatayı raporla ama sırrı yine de sil */
}
/* ... anahtarı kullan ... */
cen429_sil(anahtar, sizeof anahtar);
munlock(anahtar, sizeof anahtar);
```

`madvise(p, n, MADV_DONTDUMP)` can also be used to keep these pages out of a core dump.

```
#include <windows.h>

unsigned char anahtar[32];
if (!VirtualLock(anahtar, sizeof anahtar)) {
    /* çalışma kümesi sınırı aşılmış olabilir */
}
/* ... anahtarı kullan ... */
SecureZeroMemory(anahtar, sizeof anahtar);
VirtualUnlock(anahtar, sizeof anahtar);
```

Locking works at the **page** level (typically 4 KB). When you lock a small array, the **entire page** containing it gets locked. The amount of memory that can be locked is limited; so lock only a small, separate region that holds keys — not the whole program.

All together: a secure buffer

Gathering the rules above into a single small module makes it easier to get this right everywhere:

gizli_tampon.h — the single entry point for secrets

```
typedef struct {
    unsigned char *veri;
    size_t boyut;
} GizliTampon;

/* Ayırır, sıfırlar, belleği kilitler. Başarısızsa -1. */
int gizli_ayir(GizliTampon *t, size_t boyut);

/* Güvenli siler, kilidi açar, serbest bırakır; t->veri = NULL yapar. İki kez çağrılabilir. */
void gizli_birak(GizliTampon *t);
```

gizli_tampon.c — Linux implementation (Windows: VirtualLock/SecureZeroMemory)

```
#include <stdlib.h>
#include <sys/mman.h>
#include "gizli_tampon.h"

int gizli_ayir(GizliTampon *t, size_t boyut)
{
    t->veri = calloc(1, boyut);
    if (!t->veri) return -1;
    t->boyut = boyut;
    (void)mlock(t->veri, boyut);          /* başarısızlık ölümcül değil */
    return 0;
}

void gizli_birak(GizliTampon *t)
{
    if (!t->veri) return;                  /* çift bırakmaya karşı */
    cen429_sil(t->veri, t->boyut);
    munlock(t->veri, t->boyut);
    free(t->veri);
    t->veri = NULL;
    t->boyut = 0;
}
```

In C++ the same idea is put into a class's **destructor**; that way, once the object goes out of scope — even if an exception is thrown — the erasure happens **automatically**. This is called **RAII** (Resource Acquisition Is Initialization):

C++ — automatic erasure with RAII

```
class GizliAnahtar {
public:
    explicit GizliAnahtar(std::size_t n) : v_(n) { kilitle(v_.data(), n); }
    ~GizliAnahtar() { cen429_sil(v_.data(), v_.size()); kilidi_ac(v_.data(), v_.size()); }
    GizliAnahtar(const GizliAnahtar&) = delete;          /* kopya yok: sır çoğalmasın */
    GizliAnahtar& operator=(const GizliAnahtar&) = delete;
    unsigned char* veri() { return v_.data(); }
private:
    std::vector<unsigned char> v_;
};
```

Deleting the copy constructor is a **deliberate** decision: copying an object that holds a secret means a **second copy** somewhere else in memory that might be forgotten and never erased.

General rule in C++: no bare new/delete

Use `std::unique_ptr` for ownership (a single owner), `std::shared_ptr` if sharing is genuinely needed; use `std::vector` or `std::array` for arrays. These types free memory in their destructor; most double-free, leak, and mismatch bugs disappear on their own. That said, even `unique_ptr` frees memory **without erasing it**; secrets still need a dedicated type like the one above.

Tools that find memory errors

Tool	Platform	What it finds	How?
AddressSanitizer	GCC, Clang, MSVC	Overflow, UAF, double free	<code>-fsanitize=address /</code> <code>/fsanitize=address</code>
LeakSanitizer	GCC, Clang (Linux)	Leaks	Comes bundled with ASan; reports at program exit
MemorySanitizer	Clang (Linux)	Uninitialized reads	<code>-fsanitize=memory</code>
Valgrind Memcheck	Linux	Everything (slow, but needs no recompilation)	<code>valgrind --leak-check=full ./prog</code>
CRT debug heap	MSVC	Leaks, overflow (at the block boundary)	The <code>_CRTDBG_LEAK_CHECK_DF</code> flag via <code>_CrtSetDbgFlag</code>

How does an evaluator test this?

After a sensitive operation (login, decryption, payment) finishes, they dump the process's memory and search that dump for a known secret (a test password, a test key). If they find it, a finding is written up: "the asset persists in memory after use." In the source code, they look for a matching `free` for every `malloc`, and an erase call for every secret; they don't want to see `realloc` or by-value structs on buffers that hold secrets.

Try it yourself: memory

1. Does `char *p = malloc(10); free(p); free(p);` always crash? Why is it "dangerous even if it doesn't crash"?
2. How do you erase a password held in `std::string parola;` once you're done with it? Is `parola.clear()` enough?
3. A server leaks 100 bytes on every request. If it receives 1000 requests a second, how long does it take to exhaust 4 GB of memory? Which STRIDE letter is this?

**Answer hints**

1. No; depending on the version of the memory manager it may not even show a difference. Even when it doesn't, its internal lists get corrupted, and the next two `malloc` calls can hand the same block to two different places, leading to a situation similar to UAF. Modern glibc usually catches a simple double free and stops the program (`free(): double free detected`).
2. `clear()` only resets the length, it does not erase the bytes. Erase the content first with `cen429_sil(parola.data(), parola.capacity());` because of SSO and earlier copies, the safest option for secrets is to use a dedicated type (e.g. `GizliAnahtar`) instead of `std::string`.
3. $4 \cdot 10^9 / (100 \cdot 1000) = 40,000$ seconds $\approx$ 11 hours. **D** (denial of service). An attacker can shorten this time by increasing the request rate.

18. Protected code partitioning and secure processing under encryption

Not every line of a program is equally sensitive. Code that draws a menu on screen and code that decrypts data with a key do not deserve the same level of care. **Protected code partitioning** means dividing the program into regions by sensitivity and turning the most sensitive region into a small, isolated **core** equipped with the strongest protections. **Secure processing under encryption** means ensuring that sensitive data (even sensitive code) is **never in the clear** outside this core.

Why partition? The trusted computing base

In security, the **trusted computing base** (TCB) is all the code of which we say "if this part is buggy, security collapses." The bigger the TCB, the higher the chance it contains a bug, and the harder it is to review. The practical translation of Saltzer and Schroeder's **economy of mechanism** principle (section 4) is: **shrink the TCB**.

Example: in a 200,000-line application, if the code that works with keys is 2,000 lines and is split into a **separate module**, penetration testing and code review can **concentrate** on those 2,000 lines. Obfuscation and integrity checking are applied only to that module; the performance cost is paid only there too (see "The cost of protection").

Bölümleme: güvenilir hesaplama tabanını küçült

aralarından yalnız tutamaç ve şifreli veri geçer

GENEL BÖLGE

büyük · hızlı değişir

Arayüz
İş mantığı

hata olasılığı YÜKSEK

KORUNAN ÇEKİRDEK

küçük · seyrek değişir

Dar arayüz (tutamaçlar)
Kripto ve anahtar yönetimi
Bütünlük / RASP denetimleri

incelenebilir, kanıtlanabilir

Güvenlik kritik kod ne kadar küçükse, doğruluğunu kanıtlamak o kadar kolaydır.

Partitioning options

Where you put the protected core depends on your attacker model and your platform:

Option	Isolation strength	Example	Cost
A separate module in the same process	Low: same memory, same privilege	The C/C++ (native) library inside a Java application	Cheap; only makes the code harder to read
A separate process (privilege separation)	Medium: separate memory, different privilege	A browser's sandbox processes; a small privileged helper process (Recipe 1.4)	The design and cost of inter-process communication (IPC)
An OS service	Medium–high	Windows DPAPI, macOS Keychain, Android Keystore	Platform dependence
Hardware-backed isolation	High	Secure element, smart card, TPM, an ARM TrustZone-based trusted execution environment (TEE)	Not present on every device; hard to develop for
Moving it to the server	Highest (for the client)	Never sending the secret to the client at all, keeping it on the server, in an HSM	Requires network connectivity; doesn't work offline

The strongest protection is usually **removing the secret from the client entirely**. When that isn't possible (e.g. the phone needs to make offline payments), hardware support is sought. When that isn't available either (or isn't present on every device), the secret has to be protected **in software**, and this is exactly the situation the obfuscation, RASP, and white-box techniques we'll see through the term are for.



Partitioning in mobile payment

In a payment app that does **not** use the phone's secure element (it emulates the card in software), the typical partitioning looks like this: the user interface and workflow are in Java/Kotlin; cryptographic operations and security checks are in a **native library**. While Java bytecode can easily be decompiled back to source code, understanding compiled and obfuscated native code takes far more effort. The native layer is called **only** by the Java layer and **only** through a narrow interface (JNI); the two layers check each other's **integrity**.

The interface between partitions: narrow, validated, secret-free

The value of partitioning depends entirely on the **quality of the boundary interface**. If you keep the key in a separate module and then hand it out through a function called `anahtar_al()` ("get key"), you have gained **nothing**. Three rules for the interface:

1. **Don't hand out the secret, do the work.** The core does not return the key; it **performs the operation** with the key and returns only the result. What's handed out instead of the key is a **handle** (an opaque identifier).
2. **Validate everything that crosses the boundary.** The core treats every parameter coming from the outer region like untrusted input (size, range, format).
3. **Never pass data across the boundary in the clear.** Sensitive data that must cross the boundary does so encrypted and integrity-protected. Otherwise, the attacker hooks the interface function and reads the data **in transit**.

kasa_cekirdek.h — an interface that never hands out the secret

```
typedef struct KasaAnahtari KasaAnahtari;          /* opak tür: içi dışarıdan görünmez */

/* Ana paroladan anahtarı türetilip çekirdeğin içinde tutar; parolayı siler. */
int kasa_ac(const char *parola, size_t parola_uzunlugu,
            const unsigned char tuz[16], KasaAnahtari **cikti);

/* Anahtarla şifreler / çözer; anahtarın kendisi hiçbir zaman dışarı çıkmaz. */
int kasa_sifrele(KasaAnahtari *a, const unsigned char *acik, size_t n,
                unsigned char *sifreli, size_t kapasite, size_t *yazilan);
int kasa_coz(KasaAnahtari *a, const unsigned char *sifreli, size_t n,
            unsigned char *acik, size_t kapasite, size_t *yazilan);

/* Anahtarı güvenli siler ve serbest bırakır. */
void kasa_kapat(KasaAnahtari *a);
```

In C, an **opaque type** (a `struct` that is only declared, with its definition not in the header file) blocks calling code from directly accessing the key's bytes **at the language level**. Of course, an attacker in the same process can still read the memory; an opaque type prevents **programming mistakes**, not an **attacker**. Additional layers are needed against an attacker.

Misusing the interface: the decryption oracle and code lifting

Hiding the key well is not enough. The attacker can use **your own function** without ever finding the key:

- **Decryption oracle:** The attacker calls the `kasa_coz` function directly with data **of their own choosing**. The key stays secret, but the attacker gets to decrypt anything they want.
- **Code lifting:** The attacker **copies** the protected function (even with its embedded key) straight out of the binary and runs it as a black box **in their own program**. They never need to extract the key at all.

The countermeasures, therefore, go beyond "hide the key":

Countermeasure	Idea	Week
Binding to the device and the instance	The key is mixed with information unique to the device and to that particular install of the app; code copied into another environment does not produce the right result	3, 6
Binding to the version	Version information feeds into key derivation; the key does not resolve on an old or modified version	6
Context checking	The core verifies the caller's integrity and the expected flow; it refuses an out-of-order or out-of-context call	6
Rate and usage limits	The number of decryption operations is capped; one-time keys are used	3
No shared static key	Breaking one installation does not affect the others	3, 10

⚠ One key, every user


An application that uses the **same** embedded key in every copy loses **every single user** the moment one copy is broken. That's why certification requirements state explicitly that "breaking one instance must not affect the other instances." Keys unique to each install or device are the answer to this requirement.

Secure processing under encryption: narrowing the window of exposure

A sensitive piece of data exists in three states throughout its lifecycle: **at rest** (disk), **in transit** (network), and **in use** (memory, registers). Encrypting the first two is standard; the hard one is the **third**, because the data has to be decrypted at some point to be processed. (We will cover protecting data at rest and in transit in detail [in week 3](#).) The goal is to narrow this **window of exposure** as much as possible, in both time and space:

Açıklık penceresi: sır ne kadar süre açık?

aynı işi yapan iki tasarım, çok farklı risk



Bellek dökümü alan saldırgan için önemli olan tek şey: o anda anahtar açık mıydı?

The techniques that achieve this:

1. **Just-in-time decryption:** the data is decrypted only inside the function that will use it, and erased before that function returns.
2. **Masking in memory:** while it's not in use, the key is kept **XORed** with a random mask generated at runtime. Instead of the key itself sitting in one place in memory, there are two meaningless pieces.
3. **Splitting:** the key, or some value, is broken into pieces that sit in different places in memory (even inside different functions); they are only combined at the moment of use.
4. **Code encryption (dynamic decryption):** not just data — the code of the most sensitive functions sits encrypted in the binary too; it is decrypted into memory at runtime, run, then **re-encrypted**. Details in weeks 4 and 9.
5. **Processing without ever exposing the key:** in white-box cryptography, the key is embedded into computation tables in such a way that it never exists in the clear in memory during the operation ([week 11](#)).

Masking in memory: the key sits in two pieces

```

#include <stdint.h>
#include <stddef.h>

typedef struct {
    uint8_t maske[32];    /* açılışta rastgele üretilir */
    uint8_t ortulu[32];   /* anahtar XOR maske */
} OrtuluAnahtar;

/* Yalnız kullanım anında, çağıranın verdiği geçici tampona birleştirir. */
static void anahtari_ac(const OrtuluAnahtar *o, uint8_t gecici[32])
{
    for (size_t i = 0; i < 32; i++)
        gecici[i] = o->ortulu[i] ^ o->maske[i];
}

void imzala(const OrtuluAnahtar *o, const uint8_t *veri, size_t n, uint8_t etiket[32])
{
    uint8_t k[32];
    anahtari_ac(o, k);
    kripto_hmac_sha256(k, sizeof k, veri, n, etiket); /* işlem: code/common/cen429_kripto.h */
    cen429_sil(k, sizeof k);                          /* açıklık penceresi burada kapanır */
}

```

⚠ Masking is an obstacle, not a fortress

XOR masking prevents the key from being found by a **direct search** in a memory dump; it does not stop a determined attacker who finds the two pieces and combines them. Its value shows up when it's **combined** with other layers (obfuscation, debugger detection, frequently refreshing the mask). In security, there is no such thing as a countermeasure that is "enough on its own."

🔧 How is this applied in the field?

In a sensitive library, a single static key that must be embedded in the source code (e.g. the key used for string obfuscation) doesn't sit in one place: it is broken down at the **nibble** (half-byte) level and scattered across dozens of different functions, only combined at the moment of use. Values like a device fingerprint are stored in **two different places** with **two different hashes** and compared in the background; if either one is altered, the inconsistency shows up. Values on the application side that identify identity are also stored **XORed with the device ID**; data copied to another device becomes meaningless.

The limits of partitioning

- A module inside the same process provides **no real isolation** against an attacker who runs in that same process (e.g. a tool that hooks the program). In that case, the real protection is RASP and obfuscation.
- At an inter-process or client–server boundary, **the communication channel itself** is a new attack surface; it needs mutual authentication and integrity ([week 3](#)).
- Hardware isolation (TEE, secure element) is strong, but the code running inside it still has to be bug-free; buffer overflows have been found in TEE applications too.

**How does an evaluator test this?**

They first find the boundary: which functions are the entry points into the protected region? Then they hook those functions and log the data crossing the boundary ("is data crossing the interface in the clear?"). Next they try calling the protected function directly with inputs of their own choosing (an oracle), and copying it to another device or environment and running it there (code lifting). Finally, they dump memory during and after the operation and check whether the key sits there in the clear, or in an easily recombined form.

19. Secure development process: keeping the plan alive

A protection plan is not tested on the day it's written — it's tested every time the product **changes**. A new feature, a bug fix, or a library update can unknowingly neutralize a carefully built countermeasure. This section describes the processes that let the plan **live alongside** the product. In certified products, these processes are reviewed at least as closely as the code itself.

The secure software development lifecycle

Microsoft's Security Development Lifecycle (SDL) and NIST's Secure Software Development Framework (SSDF, SP 800-218) both say that security must be **spread across every stage** of software:

Stage	Security activity	Where it appears in this course
Requirements	Write the security requirements (derived from standards)	Week 13
Design	Threat model, protection plan, attack surface analysis	Weeks 1–2
Implementation	Coding rules, banned functions, static analysis, code review	Weeks 4–5
Verification	Sanitizers, fuzzing, penetration testing, security tests	Weeks 4, 12
Release	Signed build, unique version identity, incident response plan	This section, week 10
Response	Receiving vulnerability reports, the emergency update process	Week 12

Unique version identity: "is what was reviewed the same as what shipped?"

When an evaluator reviews a product, their report is valid for **exactly that binary**. A file shipped under the same version number but differing by even a single line is a product that was **never reviewed**. That's why every build must carry a **unique, verifiable identity**:

1. **Semantic version** (e.g. `1.4.2`): for humans.
2. **Source code identity**: the hash of the git commit the build was made from.
3. **Binary hash**: the SHA-256 of the shipped file; the user or evaluator can verify for themselves that the file they downloaded matches this value.

CMakeLists.txt — embed the git identity into the build

```
execute_process(COMMAND git rev-parse --short=12 HEAD
                WORKING_DIRECTORY ${CMAKE_SOURCE_DIR}
                OUTPUT_VARIABLE GIT_KIMLIK OUTPUT_STRIP_TRAILING_WHITESPACE)
target_compile_definitions(kasa PRIVATE SURUM="1.4.2" GIT_KIMLIK="${GIT_KIMLIK}")
```

main.c — print the version identity

```
if (argc > 1 && strcmp(argv[1], "--surum") == 0) {
    printf("kasa %s (%s)\n", SURUM, GIT_KIMLIK);
    return 0;
}
```

```
Get-FileHash .\kasa.exe -Algorithm SHA256
```

```
sha256sum ./kasa
```

**How is this applied in the field?**

The version information doesn't just stay a string printed to the screen; it also feeds into **key derivation**. That way, if a user rolls back to an old, vulnerable version (a downgrade attack), that version cannot decrypt the new version's keys. This is how version identity turns into a **security countermeasure**.


The same discipline comes back at the end of the term: in [Week 14](#), once obfuscation enters the build pipeline, which version was produced with which transformation and seed is again tracked through this identity and its hash values.

Change management: seven steps

In a mature development team, no change goes **directly to the main branch**. The typical process, adapted from IT service management, has seven steps:

Değişiklik yönetimi: yedi adım

her değişiklik izlenebilir ve geri alınabilir olmalı



The most important step from a security standpoint is the **security impact analysis** done between steps 3 and 4: "which asset, which interface, which threat, and which countermeasure does this change touch?" If the answer is "none," the change can move fast; if it touches a countermeasure, that countermeasure's tests are re-run, and the evaluator is notified if needed. In certified products, an evaluation that re-reviews **only the changed part** based on this analysis is called a **delta evaluation** (week 13).

The same discipline is applied in the source code repository:

- Main branches are **protected**; only authorized people can merge into them.
- Every change is reviewed by **at least one other developer**.
- Releases are **tagged** (preferably with a signed tag).
- External access to the repository requires **multi-factor authentication**.
- The **build environment** is protected too: the supply-chain attacks we'll see in [week 2](#) target code that gets into the build server.

The trade-off log: a record of conscious decisions

Security is almost always a **trade-off** with something: performance, usability, ease of maintenance. A mature team keeps these decisions **in writing**. A trade-off that isn't logged can't answer "why is it like this?" six months later, and looks like a **bug** in an evaluator's eyes.

Decision	Options	Chosen	Rationale	Security impact / residual risk
Logging in release builds	On / leveled / fully removed	Fully removed (via a build macro)	Silenced logging code still leaves its strings in the binary and can be turned back on	Field debugging becomes harder; only an encrypted diagnostic output was left in
Error codes	Detailed code / success-or-failure only	Success-or-failure only (an opaque value)	A detailed code tells the attacker exactly which check they tripped	A separate, server-side diagnostic channel is needed for the support team
The native library's "debuggable" flag	Off / on	On (a justified exception)	The code integrity check needs to be able to read its own memory	Compensated for with debugger detection
Password derivation time	50 ms / 250 ms / 1 s	250 ms	Balance between how long the user waits and the cost of brute-forcing	Very weak passwords are still a risk (on the residual risk list)

The third row of the table is especially instructive: one security countermeasure (integrity checking) required **loosening** another security setting (non-debuggability). A decision like this is **defensible** when it's written down with its rationale and its compensating countermeasure; when it isn't, it's a **finding**.

How protections affect the build pipeline

Countermeasures like integrity checking also change the build process. For the program to be able to check the hash of its **own** code, the correct hash value has to be embedded into the program — but embedding the hash changes the program, which changes the hash. This loop is solved with a **multi-pass build**: build once, measure the hash, embed it obfuscated into the code, rebuild; the measured region is designed so it does not include the embedded value. In the field, five-pass signed builds are used, where values like the signing certificate's hash, the package hash, and the code hash are embedded one after another. Steps like this go wrong if done by hand, so they are turned into a **build script** and documented. In [week 6](#) we will build this loop ourselves with a small example.

Measuring the cost of protections

The last piece of the protection plan is **measurement**. Saying "we added obfuscation" is not enough; you have to show the product still runs at an **acceptable speed**. A simple method: run the critical operation (e.g. a payment, opening a vault) **20 times** in both the protected and unprotected versions, and report the shortest, longest, and average time. A concrete target — such as keeping a contactless payment's total time under half a second — determines which protection can go where.

20. This week's toolbox

Bu haftanın araç kutusu

hangi araç hangi hatayı yakalar

-Wall -Wextra	şüpheli kod uyarısı	derleme
-Wsign-conversion	işaretili/işaretsiz dönüşüm	derleme
-D_FORTIFY_SOURCE=2	boyutu bilinen tampon denetimi	çalışma
-fsanitize=address	sınır dışı, use-after-free	çalışma
gdb / gcore	bellek dökümü alma	inceleme
strings / objdump -d	ikilideki metin ve kod	inceleme

Derleme zamanı ucuz, çalışma zamanı kesin, inceleme araçları saldırganın da elinde.

Tool / flag	What is it for?	Where did we see it this week?
<code>-Wall -Wextra -Wsign-conversion</code>	Compile-time warnings for suspicious code	Demo 4
<code>-D_FORTIFY_SOURCE=2 (+ -O2)</code>	Checks, at runtime, library calls that write into buffers of known size	Demo 3
<code>-fsanitize=address (ASan)</code>	Catches bugs like out-of-bounds access and use of freed memory at runtime	Demos 3, 4
<code>gdb, gcore</code>	Debugger; dumps a running process's memory to a file	Demo 2
<code>strings, objdump -d</code>	Shows the text strings and machine code inside a binary	Demo 2
<code>explicit_bzero / SecureZeroMemory</code>	Memory erasure the compiler cannot remove	Demo 2
<code>posix_spawn + absolute path + a clean environment</code>	Running another program safely	Demo 1
<code>clearenv, close_range, umask(077), setrlimit(RLIMIT_CORE), prctl(PR_SET_DUMPABLE)</code>	Cleaning up inherited state in a program's first lines	Secure start-up
<code>setresuid / setresgid + getresuid · CreateRestrictedToken</code>	Permanently dropping privilege and verifying it	Secure start-up
<code>SetDefaultDllDirectories</code>	Narrowing the DLL search order on Windows	Secure start-up
<code>snprintf, strnlen, strncpy, fgets</code>	Bounded string operations	Buffer overflows
<code>calloc, reallocarray</code>	Memory allocation that checks for multiplication overflow	Memory management
<code>mlock / VirtualLock, madvise(MADV_DONTDUMP)</code>	Preventing secrets from ending up in swap or a dump	Memory management
<code>valgrind --leak-check=full, LeakSanitizer</code>	Finding leaks and UAF	Memory management
<code>git rev-parse, sha256sum / Get-FileHash</code>	Unique version identity	Secure development process

21. Term project: this week

In the term project you will develop a C/C++ application "as if it were going through a security certification" and write a **security guide** for it (20–30 pages). This week's tasks:

- ✓ Form your team (individual or up to 4 people) and choose your topic (the topic list is in the project guide).
- ✓ Create your project repository and **project plan** (work packages, timeline, task assignment) on GitHub; get the plan approved.
- ✓ Draft the first sections of the security guide:
 - **S0** Cover page and revision history
 - **S2** Product overview – what does your application do?
 - **S3** Architecture and **interface table** (in the table format from section 6)
 - **S4** Attacker model, **STRIDE table**, and an **attack tree**
 - **S5** Asset list (draft): for every asset, its location, creation → deletion, and C/I/I+ class
- ✓ Fill in **steps 0–4** of the seven-step template from the "Worked example" section, for your own project.

S3 interface table template

ID	End A	End B	Authentication	Confidentiality / integrity	Description
A	Console app	Security DLL	...	...	...

22. On your own

These exercises are not graded; they are for reinforcing the lesson. All of them are done on top of the demos in the `code/week-01` folder, and **only on your own computer**.

? Exercise 1 – Easy: another form of the PATH trap



On Linux, change the `date` command in `rapor.c` to `ls -l rapor.c`; on Windows, change `hostname` to `whoami`. Prepare a fake `ls` (or `whoami.bat`) and run it via `PATH`. Then modify `rapor_guvenli.c` so it runs the real program by its full path.

Hint



Put the fake file in the `sahte/` folder; on Linux make it executable with `chmod +x sahte/ls`. In the fixed version, on Linux the argument array becomes `{ "ls", "-l", "rapor.c", NULL }` and the path is `/bin/ls`; on Windows the path is `GetSystemDirectoryW + \whoami.exe`. Rebuild from the command line (`.\build.ps1` or `./build.sh`) and re-run the demo.

? Exercise 2 – Easy: the effect of optimisation

In Demo 2's `CMakeLists.txt`, on the `parola_memset` line, change `MOD optimize` to `MOD korumasiz` (GCC `-O0`, MSVC `/Od`), rebuild, and run the demo. Does the `memset` version manage to erase the password this time? Why?

✓ Expected result

At `-O0`, the compiler doesn't optimise, so `memset` stays and the password gets erased. This shows that **security should not depend on the optimisation level**: release builds are always optimised, which is why `explicit_bzero` must be used.

? Exercise 3 – Medium: the exact boundary of the overflow

In Demo 3, give the `giris` program (`bin/linux/giris` or `bin\windows\giris.exe`) names that are 16, 17, 18, 19, and 20 characters long. Write down the value of the `yonetici` field for each one in a table, and explain it using the memory diagram from section 13. Why is a **16-character** name dangerous too?

🔥 Hint

`strcpy` also copies the terminating `'\0'` character. Where does a 16-character name's `'\0'` get written?

? Exercise 4 – Medium: a new privilege field

Add an `int bakiye;` (balance) field to `struct oturum`, right after `ad`. Can you push the balance above 1000 with an overflow? Then show how `giris_guvenli.c` prevents this.

? Exercise 5 – Medium: a threat model

Draw a data flow diagram for your own term project, and ask all six STRIDE letters for every flow that crosses a trust boundary. Write at least 8 threats and propose a countermeasure for each.

? Exercise 6 – Hard: read a real vulnerability report

Read a published technical write-up of Heartbleed (CVE-2014-0160). (a) Which length check was missing? (b) What does it have in common with Demo 4? (c) Which CWE category does it fall under? (d) Which lines did the fix add?

🔥 Hint

The term to search for: "out-of-bounds read" (CWE-125). The fix checks whether the length declared in the message exceeds the actual record length.

? Exercise 7 – Easy: measure secure start-up



Write a small C program on WSL/Linux; at the start of `main()`, have it print the open file descriptors (the `/proc/self/fd` folder) and the number of environment variables. Run the program normally first, then with `env -i ./prog`, then with `./prog 1>&-` (standard output closed). Then add the `tanitici_duzenle()` and `ortami_temizle()` functions from the "Secure start-up" section and observe the difference.

Hint



Count the entries in the folder with `opendir("/proc/self/fd")`. With standard output closed you won't see `printf` output; write the result to `stderr` or to a file instead. Once `tanitici_duzenle()` is added, you can see descriptor 1 bound to `/dev/null` with `ls -l /proc/<pid>/fd`.

? Exercise 8 – Easy: proof of secure erasure



In Demo 2 you showed that the `parola_explicit` version really erases the memory. Write the same program once more, in C++, keeping the password in a `std::string`. Erase it with `parola.clear()` and with `cen429_sil`, and search for it in a memory dump. In which case is the password found? Explain why, using the "Memory management and security" section.

? Exercise 9 – Medium: apply the protection plan to your own project



Fill in the seven-step template from the "Worked example" section for your own term project: **steps 0–4** complete, and initial ideas for **steps 5–7**. Have at least 6 assets in your asset table, at least 8 threats in your threat table, and a likelihood × impact score for each.

Hint



Don't forget to include the code itself, configuration files, and log files in your asset list. For every threat, answer the question "at which interface, against which asset?"; if you can't answer it, the threat is written too generally.

? Exercise 10 – Medium: dangerous function hunting



Search all `.c` files in the `code/week-01` folder for dangerous functions. For every finding, make a table of (a) which file and line, (b) whether it was placed there deliberately (for the demo), © what it was replaced with in the fixed version.

```
Get-ChildItem -Recurse -Filter *.c code\week-01 |
Select-String -Pattern 'strcpy|strcat|sprintf|gets\(|scanf'

grep -rnE 'strcpy|strcat|sprintf|gets\(|scanf' code/week-01 --include=*.c
```

? Exercise 11 – Medium: catch a leak and a UAF with a tool



In a small program, (a) call `malloc` in a loop and never `free`, (b) read a block after `free`ing it. Compile and run the program with `-fsanitize=address -g` on WSL, and with `/fsanitize=address /Zi` on Windows. For each report, read and write down the answers to "what happened, where did it happen, where was the block allocated, and where was it freed."

Hint



A UAF report has three stack traces: where the access happened, where the block was freed, and where the block was allocated. The leak report on Linux is given by LeakSanitizer as the program exits; on Windows ASan does not report leaks, so use the CRT debug heap for that.

? Exercise 12 – Medium: design an interface with an opaque handle



Looking at the `kasa_cekirdek.h` example from the "Protected code partitioning" section, write a header file for the most sensitive operation in your own project. Rule: no function returns the key or a secret in the clear; every function returns success/failure; how each parameter is validated is written as a comment.

? Exercise 13 – Hard: a masked key in memory



Get the "Masking in memory" example working: generate a random mask at start-up (Linux `getrandom`, Windows `BCryptGenRandom`), mask the key, and implement `imzala()` with the HMAC in `code/common/cen429_kripto.h`. Then dump the process's memory and search for the key: it should **not** be found while no work is in progress. If you dump memory **during** the operation (e.g. by pausing inside `imzala()` with a `getchar()`), what do you see?

? Exercise 14 – Hard: unique version identity



Embed the git identity into your own project's `CMakeLists.txt` and add a `--surum` (version) option to the program. Build at two different commits and compare the SHA-256 values of the resulting binaries. When you build the same source code twice, does the hash come out the same? If not, why not? (Look into the concept of "reproducible builds.")

Hint



Compilers can embed information like timestamps and absolute file paths into the binary. `/Brepro` in MSVC, and `-ffile-prefix-map` plus the `SOURCE_DATE_EPOCH` environment variable in GCC, reduce these differences.

23. Self-check

? 1. Explain the difference between an asset, a threat, and a vulnerability, with an example. ✓

Asset is the thing being protected (a user's password). **Threat** is the possible event/person that could harm it (someone who wants to steal the password). **Vulnerability** is the weakness the threat can exploit (the password staying in memory unerasable).

? 2. In a money transfer, changing the recipient account in transit breaks which CIA goal? Which countermeasure prevents it? ✓

Integrity. A message authentication code (MAC) or a digital signature; also TLS.

? 3. What distinguishes a 'white-box' attacker from a remote attacker? Give an example. ✓

A white-box attacker has full access to the device the program runs on: they read memory, attach a debugger, modify the binary. Example: someone analysing a mobile payment app on their own (rooted) phone.

? 4. What threat does the 'E' letter in STRIDE describe? Which demo this week is an example of it? ✓

Elevation of privilege. Demo 3: changing the `yonetici` flag with a long username.

? 5. What is the difference between an AND node and an OR node in an attack tree? Which is better from a defence standpoint? ✓

In an OR node, one branch is enough; in an AND node, every branch is required. Forcing the attacker into AND nodes (making them break several layers at once) is better for defence; that is the logic of defence in depth.

? 6. Why is calling the `date` command with `system()` dangerous? List three fixes. ✓

The shell searches for `date` via `PATH`; whoever controls `PATH` chooses which program runs. Fixes: an absolute path (`/bin/date`), running without a shell (`posix_spawn` / `execve`), a cleaned environment.

? 7. In Demo 2, why did `memset` fail to erase the password? ✓

The compiler saw that the array is never read after `memset` and removed the write as a "dead store" (dead store elimination). `explicit_bzero` is defined in a way the compiler cannot remove.

? 8. In Demo 3, why couldn't AddressSanitizer catch the 17-character attack?

ASan does bounds checking at the **object** level; because the overflow stayed inside the same struct (`ad → yonetici`), it never touched the forbidden zone. It caught it once the overflow went outside the struct. `_FORTIFY_SOURCE` , on the other hand, knew the target field's size, so it caught the short overflow too.

? 9. What happens when `int uzunluk = -1`; is passed to `memcpy`?

`memcpy` 's size parameter is `size_t` ; once `-1` is converted to unsigned, it becomes $2^{64} - 1$, and the function tries to copy a gigantic region, crashing with a memory error (or worse, corrupting memory).

? 10. What is defence in depth? Give an example from this week's mobile payment architecture.

Not relying on a single countermeasure, but building complementary layers. In the example architecture, server–phone traffic is protected both by TLS **and** by message-level encryption on top; the attacker has to break both at once.

? 11. Which of the seven protection layers are unnecessary for a server-side web service? Why?

Obfuscation (4) and RASP (5) are generally unnecessary: the attacker cannot reach the server's code or memory, they can only send input. Secure design, secure coding, compiler/OS protections, cryptography (TLS), and assurance are still needed.

? 12. Why is the 'S' (spoofing) question never asked of a data flow in a data flow diagram?

A data flow has no identity; spoofing is an **entity** (a user, a process) appearing as someone else. A data flow can be eavesdropped on (I), altered (T), or cut off (D).

? 13. List the four responses that can be given to a threat, and give an example of 'elimination.'

Mitigate, eliminate, transfer, accept. Example: never offering a "remember my password" feature entirely eliminates the threat of the password being stored on disk.

? 14. Why do you need to copy the result of `getenv("TZ")` before calling `clearenv()`?

`getenv` returns a pointer into the environment block; `clearenv` invalidates that block. Using it without copying first is a use of freed memory.

? 15. In a `setuid` program, why is the group identity changed before the user identity when dropping privilege?



Changing the group identity requires privilege. If the user identity is dropped first, the privilege needed to change the group goes away too, and the process keeps running with the privileged group.

? 16. Why is `strncpy(d, s, sizeof d)` not a safe `strcpy`?



If the source is longer than the destination, it doesn't add a `'\0'` at the end; subsequent string operations keep reading past the end of the array. The correct approach is `snprintf`, or a length check + `memcpy` + an explicit terminator.

? 17. What kind of memory bug is Heartbleed, and why did data leak without anything crashing?



An **out-of-bounds read** (CWE-125). The server only read from memory and sent it out; since it wrote nothing and the memory it read was still the process's own (valid) memory, the program never crashed.

? 18. What does the stack canary do when it notices an overflow? Why is this still a security problem?



It terminates the program immediately. Hijacking control flow is prevented, but the attacker can still crash the program whenever they want; that's a denial of service. Also, the canary doesn't see overflows that never reach the return address (e.g. an adjacent variable being modified).

? 19. What bug does the `free(p)`; `p = NULL`; habit prevent, and what does it not prevent?



It prevents double free and UAF **through that same pointer** (accessing NULL crashes immediately, and `free(NULL)` is harmless). It does nothing for **other** pointers (aliases) that also point at the same block.

? 20. Why is `realloc` dangerous for buffers that hold secrets?



If the block moves, the old location is freed **without being zeroed**; a copy of the secret stays in memory. Also, if `realloc` fails and the result is assigned directly to the same variable, the old block is lost (a leak).

? 21. What does `mlock/VirtualLock` prevent, and what does it not prevent?



It prevents the page from being written to the swap area. It does **not** prevent the memory from being read by the same process or by a debugger, the hibernation file (depending on the system), or an unerased secret; erasure is still required.

? 22. What does it mean to misuse a protected module as a 'decryption oracle'? Give a countermeasure. ✓

The attacker calls the module's decryption function directly with data of their own choosing, without ever extracting the key. Countermeasure: verify the caller's integrity and context, limit the number of calls, bind the key to the device/instance.

? 23. Using the same embedded key in every copy violates which requirement? ✓

The requirement that "breaking one instance must not affect the other instances." A key extracted from a single copy affects every user.

? 24. What are the three components of a unique version identity, and why do they matter to an evaluator? ✓

The semantic version, the source code (git) identity, and the binary's SHA-256 hash. An evaluation report is valid only for the binary that was reviewed; this is the only way to prove that the shipped file is the same one that was reviewed.

? 25. What columns does a trade-off log contain? Why is it kept in writing? ✓

Decision, options, chosen, rationale, security impact/residual risk. A trade-off that isn't written down cannot be defended later, and looks like a bug in an evaluator's eyes.

24. Resources and further reading

Textbook — J. Viega, M. Messier, *Secure Programming Cookbook for C and C++*, O'Reilly, 2003:

- Recipe 1.1 (cleaning the environment), 1.2–1.4 (restricting and separating privileges), 1.5 (file descriptors), 1.9 (preventing crash dumps)
- Recipe 3.3 (preventing buffer overflow), 3.5 (integer conversion and overflow)
- Recipe 13.2 (secure erasure from memory), 13.3 (preventing memory from being written to disk)
- Recipe 12.1 (understanding software protection)

Open sources

- J. H. Saltzer, M. D. Schroeder, "The Protection of Information in Computer Systems", 1975.
- OWASP Threat Modelling; Microsoft STRIDE.
- MITRE CWE: CWE-14, CWE-121, CWE-122, CWE-125, CWE-193, CWE-195, CWE-226, CWE-316, CWE-401, CWE-415, CWE-416, CWE-426, CWE-427, CWE-476, CWE-787.
- SEI CERT C Coding Standard: STR31-C (allocating enough space when copying strings), INT31-C (integer conversions), MSC06-C (beware of compiler optimisations), ENV33-C (do not call `system()`), MEM30-C (do not access freed memory), MEM31-C (free dynamic memory), MEM35-C (allocate sufficient memory), POS36-C (the order of dropping privilege), POS37-C (verify that privilege was successfully dropped).
- Microsoft Security Development Lifecycle (SDL); NIST SP 800-218 *Secure Software Development Framework* (SSDF).

- A. Shostack, *Threat Modelling: Designing for Security*, Wiley, 2014 (data flow diagrams and STRIDE per element).



Glossary



Term	Turkish equivalent	Short definition
Asset	Varlık	Data or resource worth protecting
Threat model	Tehdit modeli	A written definition of who the attacker is, what they can access, and what they target
Attack tree	Saldırı ağacı	An AND/OR tree showing the ways to reach an attack goal
Trust boundary	Güven sınırı	The line where components of different trust levels are separated
Buffer overflow	Arabellek taşması	Writing past the bound of a buffer
Dead store elimination	Ölü yazma giderme	The compiler removing writes whose result is never read
Sanitizer	Sanitizer	A compiler tool that monitors the program for bugs at runtime
Core dump	Çökme dökümü	A copy of a crashed process's memory written to disk
Least privilege	En az ayrıcalık	Every component running with only the privilege it needs
Defence in depth	Derinlemesine savunma	Complementary layers of protection



Next week

Week 2 – Computer viruses and security models. We will apply the threat-modelling tools we built this week (STRIDE, attack trees) and the attacker model to a concrete class of threat – malware: how a virus spreads, how it hides, and which access-control models (DAC/MAC/RBAC) stop it. The CWE classification we met for this week's memory-safety bugs will deepen next week with CVE and CVSS.