



Introduction to Secure Programming and the Application Protection Plan

CEN429 Secure Programming — Week 1

Asst. Prof. Dr. Uğur CORUH · 18.09.2026

Today's plan (3 hours)

Hour	Topic
1	Course introduction · What is security? · The attacker · Principles · Overview of application protection
2	Protection plan · STRIDE and risk score · Attack tree · Worked example "Vault" · Class exercise
3	Demo 1–2 · Secure start-up · Overflows · Demo 3–4 · Memory management · Partitioning · Secure process · Project

Learning outcomes: LO.1 (identifies and classifies vulnerabilities) · LO.5 (produces a software plan using secure design principles)

How We're Starting This Course

- This is the course's **first week**; we don't build on a previous week.
- Assumed prior knowledge: functions, arrays, and pointers in C; `cd` and `ls` in a terminal (PowerShell or Linux).
- The "**This Week's Concepts**" table below does not define this week's terms one by one — it shows **which section** teaches each one.
- The pattern we'll follow all term: **buggy code** → **attack** → **fix**.

This Week's Concepts

Each term is defined once, where it first appears in the body; here we only mark **where**.

Concept	Where
Security	Section 1
Asset	Section 1
Threat and vulnerability	Section 1
The CIA triad	Section 1
Attacker model	Section 1
White-box / MATE	Section 1
Secure design principles	Section 1
Defence in depth	Section 2
Trade-off	Section 2
Threat modelling	Section 3
STRIDE	Section 3
Attack tree	Section 3
Data flow diagram (DFD)	Section 3

How the Demos Work

- One source, two/three environments: **Windows** (Visual Studio 2022 Community or PowerShell), **WSL**, and **Linux**
- Build: Windows `.\build.ps1` · WSL/Linux `./build.sh` — build once before class
- Run: Windows `.\demo.ps1` · WSL/Linux `sh demo.sh`
- Visual Studio: **File > Open > Folder** → `code` → configuration **Windows (MSVC)** → **Build > Build All**
- 4 demos this week, under `code/week-01`

 **Ethics:** Every demo runs only inside its own folder and asks for no administrator rights. Only try these techniques on **your own computer**.

A short history — the idea of secure programming

- **1975** — Saltzer & Schroeder: the **8 principles** of secure design (least privilege, defence in depth...)
- **1970s–80s** — the **CIA triad** becomes a common language
- **1998–99** — **STRIDE** at Microsoft; Schneier introduces **attack trees**
- **2001–03** — *Building Secure Software* and the **Secure Programming Cookbook** (the course's main source)

Today's tools (CIA · attacker model · STRIDE · attack tree) are products of this lineage.

The five decisions of secure programming — diagram

Neden “güvenli programlama”? Beş karar, kodun içinde

güvenlik duvarı hatalı yazılmış programı kurtaramaz

Girdiye güvenme

uzunluk, tür, aralık, izin verilen karakterler — her seferinde

Sırrı kısa tut

işin bitince güvenle sil; bellekte bekletme

Ortama güvenme

ortam değişkenleri, dosya izinleri, çağrılan programlar

Güvenli tarafta kal

hata olursa reddet, kapat, sızdırma

Korumaları aç ve sına

derleyici bayrakları, sanitizer, fuzzing

Saldırgan izin verilen kapıdan girer: programın kendisinden.

How the course runs

- **One term project:** a C/C++ application + a **security guide** written "as if it were going through certification"
 - Midterm check: **week 7** demo · Final check: **week 15** demo
- **Two quizzes:** week 8 (1–6) · week 16 (9–14)
- Grade: $\text{Midterm} = 0.6 \cdot \text{Project1} + 0.4 \cdot \text{Quiz1}$ · $\text{Final} = 0.7 \cdot \text{Project2} + 0.3 \cdot \text{Quiz2}$
- Every week: **buggy code** → **attack** → **fix**



Ethics: Only try these techniques on **your own computer** and on the demos provided.



1. What is security?


The Vault analogy

Concept	Vault	Software
Asset	Money	Password, payment key
Threat	Thief	Someone who wants to steal the password
Vulnerability	Faulty lock	An unchecked <code>strcpy</code>
Attack	Picking the lock	Becoming administrator with a long username
Countermeasure	A new lock	Length checking
Risk	Probability × impact	"If it's found, every account is gone"

Threat · vulnerability · asset — diagram

Tehdit · zafiyet · varlık · karşı önlem

risk = olasılık $\times$ etki



Zafiyet olmadan tehdit zarar veremez; karşı önlem zafiyeti kapatır.

Security's three goals: CIA

	Question	If broken
Confidentiality	Can only authorized parties see it?	Passwords leak
Integrity	Can it be changed without authorization?	A balance, a grade changes
Availability	Does it work when needed?	The system crashes

+ Authentication · Authorization · Non-repudiation

Question: A banking app writes the balance unencrypted → ? · The recipient IBAN changes in transit → ? · It crashes on a holiday night → ?

The CIA triad — diagram

CIA üçlüsü

her güvenlik hedefi bu üçünden birine dayanır

GİZLİLİK
yalnız yetkili
GÖREBİLİR

parola sızarsa
ihlal

BÜTÜNLÜK
yalnız yetkili
DEĞİŞTİREBİLİR

bakiye değişirse
ihlal

ERİŞİLEBİLİRLİK
gerektiğinde
ÇALIŞIR

hizmet durursa
ihlal

Bir saldırıyı sınıflarken önce sorun: bu üçünden HANGİSİ bozuldu?

Answer to the question: balance unencrypted → **C**; IBAN changes in transit → **I**; crashes on a holiday night → **A**.

The CIA triad — typical countermeasures and related concepts

Goal	Typical countermeasure
Confidentiality	Encryption, access control
Integrity	Hash, MAC, digital signature
Availability	Redundancy, resource limiting

+ **Authentication**: are you really who you say you are? · **Authorization**: do you have permission to do this? · **Non-repudiation**: who did this, and can they deny it afterward?

Bug $\neq$ vulnerability, but...

A very large share of security vulnerabilities are **ordinary software bugs**.

Heartbleed (2014, CVE-2014-0160)

- Client: "Send me **64 KB** back" — but sends only **a few bytes**
- The server never **checks** the length → sends back neighboring memory (passwords, keys)
- A single missing length check → a significant part of the internet was affected

Who is the attacker, what can they reach?

Attacker	Access	Defence
Remote	Only network messages	Input validation, TLS
Local user	Files, environment variables	Permissions, secure start-up
Insider	Source code, server	Privilege separation, audit logging
Device owner (white-box)	Memory, debugger, binary	RASP, obfuscation, white-box

What's different about this course: what do we do if the program runs **on the attacker's device**?


White-box / MATE

- **MATE (Man-At-The-End)**: an attacker who **owns** the program.
- Reads the code, sees the memory, modifies it.
- The real situation for mobile/desktop applications (this course's main theme).

Attacker models — diagram

Saldırgan modelleri: ne görebiliyor?

soldan sağa görünürlük artar → savunma zorlaşır



Bu derste varsayılan model BEYAZ KUTU'dur: saldırgan cihazın sahibidir.

Attackers — real-world examples

Attacker	Example
Remote	Someone attacking a web server
Local user	Another student on a shared server
Insider	An employee abusing their role
Device owner (white-box)	Someone trying to crack a mobile payment app on their own phone

Apply each row to your own project: which one is your attacker?

Secure design principles (Saltzer & Schroeder, 1975)

Principle	One sentence
Least privilege	Run with only the privilege needed
Secure default	The default is "no permission"
Complete mediation	Check every access, every time
Open design	Security rests on the secrecy of the key , not the algorithm
Economy of mechanism	Simple = fewer bugs
Separation of privilege	Two conditions for a critical operation
Least common mechanism	Reduce shared resources
Acceptability	Unusable security gets bypassed

+ **Defence in depth**: don't rely on one countermeasure, build **layers**.

Saltzer and Schroeder's principles — diagram

Saltzer ve Schroeder (1975): güvenli tasarımın omurgası

elli yıl sonra hâlâ geçerli sekiz ilke



Bu ilkeler bir denetim listesi değil, tasarım kararlarını tartarken kullanılan ölçülerdir.

Secure design principles — examples (1/2)

Principle	Example
Least privilege	The reporting program runs as an ordinary user, not administrator
Secure default	A new user's <code>yonetici</code> (admin) field starts at 0
Complete mediation	Permission is not checked once and then cached
Open design	You do not write your own encryption algorithm

No principle is abstract; each one shows up as a concrete decision in code.

Secure design principles — examples (2/2)

Principle	Example
Economy of mechanism	A clear 20-line check beats a "clever" 500-line one
Separation of privilege	Both a password and an SMS code for a money transfer
Least common mechanism	Every session gets its own temporary file
Psychological acceptability	Asking for a password every 5 minutes leads to it being written on paper

Unusable security is **bypassed** security.

2. Overview of application protection

Defence in depth

- **Defence in depth:** not a single countermeasure, but **many**.
- If one is bypassed, another stops it.
- A "security shell" is the combination of these layers.

Why a single countermeasure is not enough

Countermeasure	What it doesn't solve
Bug-free code	The attacker opens the binary and reads the logic
Obfuscated code	Stops it with a debugger and reads the memory
Memory protected against monitoring	The key sits in code as plain text → a single string scan is enough
Hidden key	Changes a single check and bypasses the protection

→ Each layer closes the previous one's gap; the attacker must defeat **all of them at once**


Seven layers: inside out

#	Layer	Question	Week
1	Secure design	What, against whom?	1–2
2	Secure coding	Is my code bug-free?	1, 4, 5
3	Compiler/OS protections	If a bug slips through, is exploitation hard?	4
4	Obfuscation	Can someone reading the code understand it?	4, 5, 9, 14
5	RASP	Does it notice it's under attack?	6
6	Cryptography	Even if captured, is the data meaningless?	3, 10, 11
7	Assurance	How do we prove it works?	12–13

Seven layers — diagram (inside out)

Yedi koruma katmanı

aşağıdan yukarı: her katman bir öncekinin üzerine kurulur



Bir katman aşılsa bile diğerleri durdurur — derinlemesine savunma.

Even if a bug slips past the outer layers, the **next** layer tries to stop it (defence in depth).

Seven layers — typical techniques and the attacker they stop

#	Layer	Typical techniques	Attacker it stops
1	Design	Threat model, asset/interface tables, least privilege	Anyone looking for a design flaw
2	Coding	Input validation, bounds checking, SEI CERT rules	An attacker sending input remotely
3	Compiler/OS	Canary, ASLR, DEP/NX, CFI, <code>_FORTIFY_SOURCE</code>	An attacker exploiting a memory bug
4	Obfuscation	Control-flow flattening, opaque predicates, virtualisation	A reverse-engineering analyst
5	RASP	Debugger/hook/root detection, integrity checking	An attacker inspecting the running program
6	Cryptography	Authenticated encryption, key hierarchy, white-box	An attacker who captures the data
7	Assurance	Code review, fuzzing, penetration testing, certification	Auditors and evaluators

The security shell: where the layers meet

A sensitive key is carried inside **several** layers **at once**:

1. Stays **encrypted** on disk
2. **Decrypted** only at the moment it's needed, only as much as needed
3. The code section where it is decrypted is **obfuscated** and **integrity-checked**
4. The process has **debugger detection** running

→ In week 3 we'll build a four-shell design step by step

Which layer, when? By attacker model

Application	Where is the attacker?	Priority
Server-side web service	On the network; no access to code or server	1, 2, 3, 6, 7
Desktop business app	On the user's computer; the user is not the attacker	1, 2, 3, 6
Mobile payment	The device may be in the attacker's hands	All; especially 4, 5, 6
Game, licensed software	The user is the attacker	4, 5
Embedded device	Physical access	2, 3, 6 + hardware

White-box attacker: owns the program itself — copies it, opens it, modifies it, re-runs it



Obfuscation is not a substitute for correct code

- Obfuscation **slows the attacker down**, it does not **fix the bug**
- An obfuscated overflow is still an overflow
- Outer layers only make sense once the inner layers are solid

The server never trusts the client. Even an obfuscated client can send inconsistent data to the server.

The white-box attacker's six paths (1/2)

A six-row **attack table** for a sensitive library running on the client side:

#	Attacker's path	Example	Layer that answers
1	Bypassing platform controls	A rooted device, an emulator, a modified OS	5 (RASP)
2	Reverse-engineering the code	Reading logic/keys with a decompiler	4 (obfuscation), 6
3	Modifying the code	Flipping an <code>if</code> check and repackaging	5 (integrity), 4

The white-box attacker's six paths (2/2)

#	Attacker's path	Example	Layer that answers
4	Abusing interfaces	Calling the library from your own program	1 (design), 5, 6
5	Extracting assets at runtime	Memory dump, debugger, hook	2 (erasure), 5, 6 (white-box)
6	Redirecting control flow at runtime	Skipping a branch, changing a return value	5 (flow counter), 4

This table is the starting point for the **threats** section of a protection plan.

Trade-off

- Every protection carries a **cost**: speed, complexity, expense.
- Security is not "infinite protection," it is **balanced** protection.
- We write the remaining risk down explicitly.

The cost of protection

Cost	Example
Performance	Control-flow flattening can slow a function several times over
Size	Virtualisation, white-box tables make the binary bigger
Maintenance	Crash reports from obfuscated code are unreadable
False positive	An overly sensitive root check punishes a legitimate user

→ Every countermeasure is written **with its justification**: *this asset, against this attacker, at this cost*

An evaluator reads the table backward

1. String scan → is a key, URL, error message visible?
2. Opening the binary → is the logic readable?
3. Attaching a debugger → does the program notice?
4. Modifying the code → does the integrity check catch it?
5. Memory dump → is a secret found?

Every step that **isn't** stopped becomes a finding in the report

Try it yourself (3 min)

1. What is the attacker model for an **offline password manager**? Which layers matter?
2. "It runs on the server, no need to obfuscate" — when is this true, when is it false?
3. In a game that uses only obfuscation, which attack path stays open?

Try it yourself — answers

1. **Offline password manager**: the attacker owns the device/file (white-box). Layers that matter: **cryptography** (a KDF from the password + AEAD), **secure coding**, some obfuscation/RASP.
2. **"Runs on the server"**: **true** if no sensitive logic/key is sent to the client; **false** if code/a key runs on the client (white-box).
3. **Obfuscation only**: if there's no server-side validation, a patched-client cheat stays open — obfuscation doesn't **fix** the logic.

3. The application protection plan and threat modelling

Threat modelling

- **Threat modelling:** systematically answering "what will we protect, against whom, and how?"
- It lists assets, threats, and countermeasures.
- Today we'll do it with STRIDE and attack trees.

Application protection plan: 7 steps

1. **Scope** — What are we protecting? Which binary, which version, which hash?
 2. **Architecture and interfaces** — Who talks to whom, over which channel?
 3. **Assets** — Where does it live, when is it erased, **C** or **I**?
 4. **Threats** — Who is the attacker, by which path do they reach the asset?
 5. **Countermeasures** — Which layer, which countermeasure?
 6. **Verification** — How is it proven that it works?
 7. **Residual risk** — What was accepted, and why
- In certification this plan becomes **a security guide hundreds of pages long**. You will write a **20–30 page** version of it for the project.

Step 1 detail: Scope

Question: What are we evaluating?

- Saying "**version 2.1**" alone is not enough
- Which **binary**, which **source code**, which **hash** value?
- In certification this is called the **target of evaluation**

Step 2 detail: Architecture and interfaces

Question: What are the components, how do they talk?

- What **are** the components?
- Which **channel** do they talk to each other over?
- How is identity **verified** on each channel?

Step 3 detail: Assets

Question: Where and when is the protected data?

- Every piece of data to be protected is listed **one by one**
- **Where** does it live, **when** is it created, **when** is it erased?
- Does it need **confidentiality (C)** or **integrity (I)**?

Step 4 detail: Threats

Question: Who is the attacker, how do they reach it?

- **Who** is the attacker? (see attacker models)
- By **which path** can they reach each asset?
- STRIDE and the attack tree come in here

Step 5 detail: Countermeasures

Question: Which countermeasure, in which layer?

- **One countermeasure** is chosen for each threat
- Which of the seven layers does the countermeasure **belong to**?
- An unjustified countermeasure is **in the wrong place** or **unnecessary**

Step 6 detail: Verification

Question: How is it proven that it works?

- A **test** is written for each countermeasure
- A unit test, or a **bypass attempt**?
- A countermeasure with no test **does not exist** in an evaluator's eyes

Step 7 detail: Residual risk


Question: What did we knowingly accept?

- Risks that can't be closed, or are knowingly **accepted**
- The **reason** for each (e.g. performance, usability)
- The plan **never ends**: as the product changes, step 1 is revisited and it is **updated**

The protection plan — diagram


Uygulama koruma planı: yedi adım

1–3 TANIMLA · 4–7 KORU ve KANITLA



Bu yedi adım dönem projesinin de iskeletidir (S2–S17).

Example architecture: a mobile payment app



- The phone may be **in the attacker's hands**: root, debugger, memory dump
- The most sensitive work runs in the **native C/C++** layer

Interface and asset tables

Interface	Ends	Authentication	Confidentiality / integrity
C	native – database	—	Encryption + MAC, key bound to the device
D	SDK – server	Certificate pinning	TLS + message-level encryption on top

Asset	Where	Lifetime	Protection
One-time payment key	DB (encrypted), memory	One payment	C, I
Session key	Memory only	Session	C, I
Transaction counter	DB	Persistent	I

An asset not on the list is an asset that's not protected.

Interface table — full list (mobile payment)

Interface	Ends	Authentication	Confidentiality / integrity
A	App – Library (Java)	The caller's signature is verified	Same process; no sensitive data crosses
B	Library (Java) – native	Mutual integrity check	Crosses encrypted and masked
C	native – local database	—	Encryption + MAC, key bound to the device
D	Library – back-end server	Certificate pinning + server verification	TLS + message-level encryption on top
E	Library – payment terminal	Payment protocol	One-time payment key

Asset table — full list (mobile payment)

Asset	Where	Creation → deletion	Protection
One-time payment key	DB (encrypted), memory	Downloaded from server → used → erased	C, I
Device fingerprint	Memory, two separate pieces	Computed at start-up	I
Session key	Memory only	Derived → erased at session end	C, I
Transaction counter	Database	Increments on every payment	I
Server public key	Embedded in the app	At build time	I

STRIDE

Threat		Breaks	Example
S	Spoofing	Authentication	A fake app calling the SDK
T	Tampering	Integrity	A native check altered in the binary
R	Repudiation	Non-repudiation	"I didn't make this payment"
I	Information disclosure	Confidentiality	Key read from a memory dump
D	Denial of service	Availability	Thousands of fake registration requests
E	Elevation of privilege	Authorization	Ordinary user → administrator

Draw the data flow diagram → ask the six letters for **every arrow that crosses a trust boundary**.

Example system

A mobile banking application:

- user login
- balance display
- money transfer

Let's apply every STRIDE letter to this system.

S · Spoofing

- **Threat:** acting as someone else (a fake user/server).
- **Example:** a fake server, a stolen session.
- **Countermeasure:** strong authentication, TLS + certificate checking.

T · Tampering

- **Threat:** modifying data/code without authorization.
- **Example:** changing the transfer amount.
- **Countermeasure:** integrity (MAC/signature), integrity checking (RASP).

R · Repudiation

- **Threat:** denying what you did.
- **Example:** "I didn't make that transfer."
- **Countermeasure:** secure log, signature, audit trail.

I · Information disclosure

- **Threat:** confidential data leaking.
- **Example:** a balance/password leak.
- **Countermeasure:** encryption, least privilege, leak-free error messages.

D · Denial of service

- **Threat:** making the service unavailable.
- **Example:** excessive requests, resource exhaustion.
- **Countermeasure:** rate limiting, resource quotas, input limits.

E · Elevation of privilege

- **Threat:** gaining unauthorized privilege.
- **Example:** an ordinary user becomes administrator.
- **Countermeasure:** least privilege, privilege checking, secure default.


STRIDE Summary for the Banking Example

Letter	Countermeasure in this system
S	TLS + authentication
T	MAC/signature + RASP
R	Secure log
I	Encryption + least privilege
D	Rate limiting
E	Privilege checking

Root · goal

Saldırı ağacı: Ödeme anahtarını ele geçir

VEYA = tek dal yeter · VE = tüm dallar gerekir



The attacker's ultimate goal. Now let's break down the paths.

Attack tree: "Capture the payment key"

GOAL: Payment key — *OR* (one is enough)

1. **Break the database** — *AND* (all required): gain root privilege + find the database key
2. **Read from memory** — *OR*: attach a debugger | take a memory dump | hook a function
3. **Eavesdrop on the traffic** — *AND*: break TLS + break message-level encryption

Attack tree — the "Break the database" branch

AND node (all required):

1. Gain root privilege
2. **AND** find the database key

This path is **expensive** because both are needed.

This is proof that defence in depth **works**.

Attack tree — the "Read from memory" branch

OR node (one is enough): three alternatives

1. Attach a debugger
2. **OR** take a memory dump
3. **OR** hook a function

Three alternatives = **the weakest point**: the key in memory.

→ This week's Demo 2 · Week 6 runtime protection · Week 11 white-box

Attack tree — the "Eavesdrop on traffic" branch

AND node (all required):

1. Break TLS
2. **AND** break message-level encryption

Both layers must be broken **at once** → an expensive path.

The defender's goal: force the attacker into as many **AND** nodes as possible.

Attack tree — mapping to countermeasures

Let's map the attack tree's leaves directly onto a defence layer:

Leaf	Countermeasure
Debugger	Anti-debug (6)
Memory dump	Short lifetime + protection (6, 10)
Extract from binary	Obfuscation + white-box (9, 11)
Leak over the network	TLS + pinning (10)

Data flow diagram: five symbols


Symbol	Name	Example
□ Rectangle	External entity	User, payment terminal
○ Circle	Process	Login module, crypto library
= Two lines	Data store	Database, config file
→ Arrow	Data flow	HTTP request, function call
--- Dashed line	Trust boundary	Internet ↔ server, app ↔ OS

➔ Color every arrow that crosses a trust boundary **red**: start your review there

Data flow diagram and trust boundary — diagram

Veri akış diyagramı: dört öge ve güven sınırı

tehdit modellemenin çizim dili



Sınırı geçen her ok bir tehdit adayıdır: orada STRIDE'in altı harfini tek tek sorun.

Kural: güven sınırını çizmeden tehdit listesi çıkarılmaz.

Which letter for which element?

Element	S	T	R	I	D	E
External entity	✓		✓			
Process	✓	✓	✓	✓	✓	✓
Data store		✓	✓ ¹	✓	✓	
Data flow		✓		✓	✓	

¹ If it's an audit log · A data flow has no "identity" → **S is not asked**

Question and countermeasure for each letter

Letter	Ask yourself	Typical countermeasure
S	Is the other party really who they claim to be?	Authentication, signature, certificate pinning
T	Would I notice if it were changed?	MAC/HMAC, signature, integrity check
R	If they say "I didn't do it," do I have proof?	Tamper-resistant log, signed transaction
I	Does it leak on disk, network, memory, or in a log?	Encryption, memory wiping, masking
D	Can it be crashed, exhausted?	Input limits, rate limiting, quotas
E	Can a low-privileged actor get a high-privileged one's work done?	Least privilege, complete mediation, memory safety

An overflow = **D** (crashes) + **E** (control flow hijacked) + **I** (adjacent secret is read)

Risk score: likelihood × impact

	Impact 1	Impact 2	Impact 3
Likelihood 3	3	6	9 critical
Likelihood 2	2	4	6
Likelihood 1	1	2	3

- **Likelihood:** physical access, or over the internet? Expertise needed? Can it be automated?
- **Impact:** which asset, how many users, financial/legal consequence?
- More detailed: **CVSS** (week 2), **attack potential** (week 13)

Four responses to a threat

1. **Mitigate** — add a countermeasure (most common)
2. **Eliminate** — never write the feature at all (strongest)
If there's no "remember my password" feature, there's no disk-password threat either
3. **Transfer** — don't store card data, leave it to the payment provider
4. **Accept** — write it into the **residual risk** list, with a reason

⚠️ A risk not written into the plan is not accepted, it has been **overlooked**

4. Worked example: the "Vault" password vault

Step 0–1: Application and scope

Vault: a C++ desktop password vault

- Master password → derived key → records encrypted in a single file
- Copy to clipboard · optional cloud backup · auto-update

Field	Value
Target of evaluation	<code>kasa</code> 1.0 + <code>kasa_cekirdek</code> + configuration
Out of scope	The backup server itself, the operating system
Security objective	Unreadable and undetectably unmodifiable to anyone who doesn't know the master password


Step 2: Architecture and interfaces

Interface	Ends	Authentication	Confidentiality / integrity
A	User → Interface	Master password	Never shown on screen
B	Interface → Core	Same process	Password erased after the call
C	Core → Vault file	—	AES-GCM, key from the password
D	Core → Clipboard	—	Cleared after 30 s
E	Core → Backup server	Certificate + token	TLS + already-encrypted file
F	Interface → Update	Package signature	TLS; the real assurance is the signature

Vault architecture — diagram

Parola kasası: bileşenler ve arayüzler

her ok bir arayüz — her arayüz doğrulanır



Step 3: Assets

#	Asset	Where	Creation → deletion	Protection
V1	Master password	Memory only	Typed → derived → erased immediately	C
V2	Vault key	Memory only	Derived → erased on lock	C, I
V3	Decrypted record	Memory	Displayed → erased on close	C
V4	Vault file	Disk, backup	Persistent	C, I
V5	Salt, derivation parameter	File header	At creation	I
V6	Update public key	Embedded in the app	At build time	I

Protection classes: **C** confidentiality · **I** integrity · **I+** integrity + accountability · **N** none

Step 4: Threats

#	Where	Letter	Threat	Risk
T1	C / V4	I	Stolen file, weak password brute-forced offline	9
T2	B / V1-2	I	Password/key stays in memory → dump, swap	6
T3	D / V3	I	Clipboard password read by another app	6
T5	C / V5	T	Iteration count dropped to 1	3
T6	F / V6	T, E	A fake update is installed	6
T8	Process	T	A memory bug in the vault parser	6

Step 3-4 (continued): remaining assets and threats

#	Asset/Threat	Detail	Value
V7	Backup token	Disk (OS key store); created on connect → deleted on revoke	C
T4	C / V4, T	Vault file is modified; a corrupted record is used without notice	Risk 4
T7	E, S	A fake backup server steals the token	Risk 4
T9	A, S	Master password read off the screen (shoulder surfing)	Risk 3

All nine threats must be in the table; a skipped threat is a threat that isn't closed.

Step 5: Countermeasures

Threat	Countermeasure	Layer
T1	Slow derivation (Argon2id / high-iteration PBKDF2) + password strength requirement	6, 2
T2	Secure erase, <code>mlock</code> , disable dumps	2
T3	Clipboard cleared after 30 s	1
T4-5	Authenticated encryption; header = AAD ; minimum iteration count hard-coded	6, 2
T6	Never run without verifying the signature; reject downgrades	6
T8	Length validation + fuzzing + compiler protections	2, 3, 7

Step 5 (continued): T7 and T9 countermeasures

Threat	Countermeasure	Layer
T7	Certificate validation + hostname check, optional certificate pinning	6
T9	Accept: mask the password on screen; physical observation is out of scope	—

The T9 example matters: not every threat needs a countermeasure **added**, sometimes **accepting** it is the right call.

Step 6–7: Verification and residual risk

Countermeasure	Test	Expected
Derivation	Time it	≥ 250 ms
Erasure	Search a memory dump after lock for the password	Not found
Integrity	Modify one byte of the file	Says "corrupted," refuses to open, doesn't crash
Signature	A package with a broken signature	Rejected, logged

Residual risk: malware running with the same user while the vault is open · shoulder surfing · a very weak password



A countermeasure with no test **doesn't exist** in an evaluator's eyes

Step 6 (continued): T5 and T8 tests

Countermeasure	Test	Expected
T5 parameter	Iteration count in the header set to 1	The app rejects the file
T8 parser	24 hours of fuzzing (week 4)	No crashes; no sanitizer findings

Step 7 (continued): residual risk — full table

Risk	Why accepted?	Re-evaluation
Memory can be read if malware runs while the vault is open	Full protection against same-user malware isn't possible; the lock timeout was kept short	Every release
Shoulder surfing (T9)	The physical environment is out of scope	—
Very weak master password	The user is warned but not forcibly blocked	Based on user feedback

Common mistakes

- **Forgetting assets:** keys, tokens, counters, configuration, **the code itself**
- **Drawing the trust boundary wrong:** treating your own client as "trusted"
- **Not matching a countermeasure to a threat:** "we use TLS" — which threat does it close, which doesn't it?
- **A countermeasure with no test:** saying "we erase it" without noticing the compiler removed the erasure (Demo 2)




5. Program start-up and secrets left in memory

How a secret leaks from memory — diagram

Sır bellekten nasıl dışarı sızar?

değişken kapsam dışına çıkar, baytlar kalır



Çözüm: `explicit_bzero` / `SecureZeroMemory` — derleyicinin kaldıramayacağı silme.

A program doesn't start with a blank page

What it **inherits** from whoever launched it:

- Environment variables: `PATH`, `LD_PRELOAD`, `IFS` ...
- Open file descriptors, working directory, `umask`
- Resource limits, crash dump settings

If the program trusts these, **whoever controls them controls the program.**

Book: Viega & Messier, Recipes 1.1–1.9

Demo 1 — Fooling via PATH (CWE-426)

```
int durum = system(KOMUT); /* Linux: "date"  Windows: "hostname" */  
                          /* HANGİSİ çalışacak? */
```

`.\demo.ps1` (Windows) · `sh demo.sh` (WSL / Linux)

```
ADIM 2 – PATH="$PWD/sahte:$PATH" ./bin/linux/rapor
```

```
Rapor tarihi:
```


```
!!! SAHTE 'date' calisti !!!
```

```
ADIM 3 – ./bin/linux/rapor_guvenli -> Rapor tarihi: Sat Sep 19 ...
```

Fooling via PATH — diagram

PATH ile kandırma: neden system() tehlikeli?

kabuk, programı adına göre PATH sırasıyla arar



Düzeltilme: mutlak yol (/bin/date) · kabuksuz çalıştırma (execve) · temiz ortam.

Fix: three rules

```
char *const arguman[] = { "date", NULL };  
char *const temiz_ortam[] = { "PATH=/usr/bin:/bin", "LANG=C", NULL };  
posix_spawn(&pid, "/bin/date", NULL, NULL, arguman, temiz_ortam);
```

1. **Absolute path** → no PATH search
2. **No shell** (`posix_spawn` / `execve`) → no special characters, no `IFS`
3. **Clean environment** → a known, small environment

Windows: `CreateProcessW` + `System32\hostname.exe` + only `SystemRoot`

Best: avoid needing another program at all → `strftime()`, `GetComputerNameW()`

Secure start-up: what is inherited?

Inherited	Danger	Recipe
Environment variables	<code>PATH</code> , <code>LD_PRELOAD</code> change the program/library that loads	1.1
Open file descriptors	The parent's hidden file leaks; if 0–2 are closed, the first file becomes "stdout"	1.5
<code>umask</code>	<code>000</code> → world-writable files	2.7
Resource limits	Core dump enabled → memory to disk	1.9
Identity and privileges	setuid: more privilege than needed	1.3

Rule: trust nothing you didn't set up yourself

What a program inherits — diagram

Program başlarken ebeveyninden ne devralır?

`main()`'in ilk satırlarında yapılacak iş budur

Ortam değişkenleri

PATH, LD_PRELOAD, LD_LIBRARY_PATH — hangi kod yüklenecek

Açık dosya tanıtıcıları

0/1/2 kapalıysa açtığınız dosya stdin olur

umask ve çalışma klasörü

oluşturduğunuz dosyanın izinleri buradan gelir

Sinyal düzeni ve kaynak sınırları

yok sayılan sinyal, düşük core limiti

Gerçek ve etkin kullanıcı kimliği

setuid ise ayrıcalık elinizdedir

Hiçbirine güvenmeyin: gerekeni kendiniz kurun, gerekmeyeni kapatın.

Other inherited risks: working directory, signals, dlopen

- **Working directory:** relative paths (`./ayar.ini`) can point into a folder the attacker chose
- **Signal settings:** signals the parent ignores are **also ignored in the child**
- `dlopen("lib.so")` : affected by `LD_LIBRARY_PATH` and search paths → use the full path; `LD_*` should already be cleaned in Step 1

Same rule: **trust nothing the program itself did not set up.**

Steps 1–2: Environment and descriptors

```
/* Ortam: kara liste değil BEYAZ liste */
clearenv();                               /* önce korunacakları kopyala! */
setenv("PATH", "/usr/bin:/bin", 1);

/* 0-2 kapalıysa /dev/null'a bağla, gerisini kapat */
for (int fd = 0; fd <= 2; fd++)
    if (fcntl(fd, F_GETFD) == -1 && open("/dev/null", O_RDWR) != fd) abort();
close_range(3, ~0U, 0);                   /* Linux 5.9+ */
```

- The `getenv` pointer becomes **invalid** after `clearenv`
- Always use `O_CLOEXEC` on your own files

Steps 3–4: Permissions and crash dumps

```
umask(077);                                /* rw----- */
int fd = open(yol, O_WRONLY | O_CREAT | O_EXCL, 0600);

struct rlimit r = { 0, 0 };
setrlimit(RLIMIT_CORE, &r);                /* döküm boyutu 0 */
prctl(PR_SET_DUMPABLE, 0, 0, 0, 0);        /* + ptrace ile bağlanılamaz */
```

- Windows: `SetErrorMode(SEM_FAILCRITICALERRORS | SEM_NOGPFAULTERRORBOX)`
- The real defence: even if a dump is taken, no secret should be **left inside it**

Step 5: Drop privilege permanently and verify

```
if (setresgid(g, g, g) != 0) abort(); /* 1. önce GRUP */
if (setresuid(u, u, u) != 0) abort(); /* 2. sonra kullanıcı: r/e/s üçü de */
getresuid(&r, &e, &s);
if (r != u || e != u || s != u) abort(); /* 3. DOĞRULA */
```

- `seteuid()` is temporary: the saved identity stays privileged (week 2 Demo 13)
- Windows: a **restricted token** (`CreateRestrictedToken`)
- Better: **privilege separation** — a small privileged process + a large unprivileged one (Recipe 1.4)

Step 6: Loading programs and libraries

✗ Wrong	✓ Correct
<code>system("convert " + dosya)</code>	<code>execve("/usr/bin/convert", argv, temiz_ortam)</code>
<code>execvp("convert", ...)</code>	<code>execve</code> with a full path
<code>CreateProcess(NULL, "C:\Program Files\...")</code>	<code>lpApplicationName</code> = full path, quoted command line
<code>LoadLibrary("yardimci.dll")</code>	<code>SetDefaultDllDirectories(LOAD_LIBRARY_SEARCH_DEFAULT_DIRS)</code> + full path

The Windows DLL search order = the twin of the `PATH` problem (CWE-427)

All together: `guvenli_baslat()`

```
int main(int argc, char **argv)
{
    tanitici_duzenle();    /* 1. tanıtıcılar */
    ortami_temizle();     /* 2. ortam      */
    umask(077);           /* 3. izinler  */
    dokumu_kapat();       /* 4. döküm    */
    yetkiyi_birak();      /* 5. yetki    */
    /* ... asıl iş ... */
}
```

Order matters: before doing anything else

Demo 2 — Password left in memory (CWE-14)

```
char parola[64];  
n = dosya_oku(fd, parola, sizeof(parola) - 1);  
anahtar = anahtar_turet(parola, n);  
memset(parola, 0, sizeof(parola));           /* sürüm 2 */  
explicit_bzero(parola, sizeof(parola));       /* sürüm 3, Linux */  
SecureZeroMemory(parola, sizeof(parola));    /* sürüm 3, Windows */
```

```
Hiç silmeyen   -> parola dökümde BULUNDU  
memset (-02)   -> parola dökümde BULUNDU    ← !!!  
explicit_bzero -> parola dökümde bulunamadı
```

Why did the memset disappear?

The compiler's reasoning:

" `parola` is never **read** after this point. Writing zeros into it doesn't change the result. **Redundant — I'll remove it.**"

Dead store elimination: good for speed, a disaster for secrets.

Platform	Erasure that can't be removed
Linux (glibc)	<code>explicit_bzero</code>
Windows	<code>SecureZeroMemory</code>
C11 Annex K	<code>memset_s</code>
OpenSSL	<code>OPENSSL_cleanse</code>

+ `setrlimit(RLIMIT_CORE, 0)` · `mlock` · keep the secret **short-lived**, don't **copy** it




6. Process memory and overflows

Process memory

Süreç belleği: verileriniz nerede duruyor?

yığın aşağı, öbek yukarı büyür — ortada boşluk



Taşmanın sonucu hangi bölgede olduğuna bağlıdır: yığında dönüş adresi, öbekte komşu blok.

C does not check the end of an array. A 17th byte into `ad[16]` → overwrites the next variable.

The structure in memory

```
struct oturum { char ad[16]; int yonetici; };
```

```
offset:  0 ..... 15 | 16  17  18  19
         [ a  y  s  e \0 . . . . ] [ 00 00 00 00 ]  yonetici = 0

"AAAAAAAAAAAAAAAAAB" (17 karakter) kopyalanınca:
         [ A  A  A  A  A  A  A  A  A  A  A ] [ 42 00 00 00 ]  yonetici = 66 (!)
                                   'B'  '\0'
```

Little-endian: the lowest byte sits at the lowest address.

Structure in memory and overflow — diagram

Yapının bellekteki hâli: taşma neyi ezer?

ad[16] ve yönetici yan yana durur

ÖNCE

"ayse"
ad[16]

00 00 00 00
yönetici = 0

SONRA

"AAAAAAAAAAAAAAAAAAAA"
17 karakterlik "AAAA...AAAB" kopyalandı

B 00 00 00
yönetici $\neq$ 0

Tek bayt taşma yetkiyi değiştirdi: sınır denetimi olmayan kopya, komşu alanı ezer.

What is a buffer overflow?

Writing **more data than allocated** into an array → the overflowing bytes land **over neighboring variables**

Why does C overflow so easily?

1. An array **doesn't know its own size**
2. Only the **starting address** goes to the function
3. A string = an array ending in `'\0'`; `strcpy` never **asks** the target's size
4. The compiler **adds no** bounds checking


CWE-787 "out-of-bounds write": always near the top of the CWE Top 25

A stack overflow

```
void selamla(const char *ad) {  
    int yetkili = 0;  
    char tampon[16];  
    strcpy(tampon, ad);          /* ad > 15 karakter? */  
}
```

Yığında taşma adım adım

strcpy aşağıdan yukarı yazar



Kanarya, dönüş adresinden önce durur: ezilirse program sonlanır (ama yetkili zaten ezilmiştir).

Types of overflow

Type	Typical cause	CWE
Stack overflow	<code>strcpy</code> , <code>gets</code> , unbounded <code>scanf("%s")</code>	121
Heap overflow	A <code>memcpy</code> with the wrong length	122
Off-by-one	<code><=</code> instead of <code><</code> ; no room for <code>'\0'</code>	193
Out-of-bounds read	Trusting a length the other party reports	125
Integer-driven	<code>n * size</code> overflows → small block, large copy	190 → 787

Heartbleed (2014): nothing was written

- Client: "I'm sending 1 byte, but reporting **64 KB**"
- Server: sent 64 KB back **without checking** the declared length
- Leaked: passwords, sessions, even the server's **private key**
- CVE-2014-0160 · CWE-125 · Fix: **a single length check**

➔ An out-of-bounds read is as dangerous as a write; it **doesn't crash, leaves no trace**

The off-by-one bug

```
char ad[8];  
for (int i = 0; i <= 8; i++)      /* 9 kez: ad[8] dışarıda */  
    ad[i] = kaynak[i];  
  
strncpy(ad, kaynak, sizeof ad); /* kaynak >= 8 → '\\0' YOK */  
printf("%s\\n", ad);             /* dizinin sonundan okumaya devam */
```

- `strncpy` is **not** a safe `strcpy`
- Even a single-byte overflow changes the low byte of a neighboring variable

Dangerous functions → use instead

✗	✓
<code>gets(s)</code>	<code>fgets(s, sizeof s, stdin)</code> — <code>gets</code> removed in C11
<code>strcpy(d, s)</code>	Length check + <code>memcpy</code> , <code>snprintf</code> , <code>strncpy</code> (glibc 2.38+)
<code>strcat(d, s)</code>	Compute remaining room; <code>snprintf</code> , <code>strlcat</code>
<code>sprintf(d, ...)</code>	<code>snprintf(d, sizeof d, ...)</code> + check the return value
<code>scanf("%s", s)</code>	<code>scanf("%15s", s)</code>
<code>memcpy(d, s, n)</code>	First check <code>n <= sizeof d</code>
<code>char a[n]</code> (VLA), <code>alloca</code>	A fixed upper bound, or <code>malloc</code>

Today's Recipe 3.3–3.4: `snprintf` / `strncpy`, MSVC `strcpy_s`, C++ `std::string` / `span`

The right pattern: validate first, then copy with a bound

```
int selamla(const char *ad)
{
    char tampon[16];
    size_t n = strlen(ad, sizeof tampon); /* en fazla 16 bayt oku */
    if (n == sizeof tampon)               /* uzunsa REDDET */
        return -1;
    memcpy(tampon, ad, n);
    tampon[n] = '\0';                     /* sonlandırıcı her zaman */
    printf("Merhaba %s\n", tampon);
    return 0;
}
```

Why reject instead of truncate? `/home/ayse/gizli_rapor.txt` → `/home/ayse/gizli`

C++: the type carries the size — but be careful

```
void selamla(const std::string &ad) {  
    if (ad.size() > 15) throw std::invalid_argument("ad cok uzun");  
    std::cout << "Merhaba " << ad << '\n';  
}
```

- `vector::operator[]` does **not** check bounds → `at()` does
- Writing through a pointer taken with `c_str()` → all of C's problems come back

Who catches the overflow?

Layer	Tool	When?
Code	Bounds checking, safe API	Always
Warning	<code>-Wall -Wextra -Wformat-security · /W4 /sd1</code>	Development
Static analysis	<code>clang --analyze</code> , <code>cppcheck</code> , <code>/analyze</code>	CI
Sanitizer	<code>-fsanitize=address · /fsanitize=address</code>	Testing
Fuzzing	libFuzzer, AFL++	Testing (week 4)
Compiler	Canary <code>-fstack-protector-strong · /GS</code> , <code>_FORTIFY_SOURCE</code>	Release
OS / hardware	DEP/NX, ASLR · Intel CET, ARM PAC	Release

⚠ Protection makes exploitation **harder**, it doesn't **fix the bug** — a canary terminates the program = denial of service

Reading an ASan report

```
ERROR: AddressSanitizer: stack-buffer-overflow
WRITE of size 21 at 0x7ffd... thread T0
    #0 in strcpy
    #1 in selamla ornek.c:7
    #2 in main      ornek.c:15
This frame has 1 object(s):
    [32, 48) 'tampon' <== Memory access at offset 48 overflows this variable
```

1. **What?** a stack overflow, a 21-byte write
2. **Where?** `ornek.c:7`, `strcpy`
3. **Which variable?** `tampon` (16 bytes), the access is right past the bound

Roughly 2× slower → only in **test** builds

Demo 3 — Privilege escalation via overflow (CWE-121)

`.\demo.ps1` (Windows) · `sh demo.sh` (WSL / Linux)

Version	17-character attack
<code>giris</code> (unprotected)	✗ ADMIN access granted
<code>giris_asan</code> (ASan)	✗ stayed silent — overflow inside the struct
<code>giris_asan</code> + 40 characters	✓ stack-buffer-overflow caught
<code>giris_denetimli</code> (<code>_FORTIFY_SOURCE=2</code> / <code>strcpy_s</code>)	✓ program stopped
<code>giris_guvenli</code>	✓ Rejected

Fix: validate, then copy with a bound

```
/* İzin listesi: 1-15 karakter, harf/rakam/_/- */
static int ad_gecerli_mi(const char *s)
{
    size_t n = strlen(s, 16);
    if (n == 0 || n >= 16) return 0;
    for (size_t i = 0; i < n; i++) {
        unsigned char c = (unsigned char)s[i];
        if (!isalnum(c) && c != '_' && c != '-') return 0;
    }
    return 1;
}

struct oturum o = { .yonetici = 0 }; /* güvenli varsayılan */
snprintf(o.ad, sizeof(o.ad), "%s", argv[1]); /* boyutu bilen kopya */
```

Canary, NX, ASLR (week 4) **make exploitation harder, they don't fix the bug.**

Integers: when does -1 become huge?

`int` is signed · `size_t` is unsigned · negatives are **two's complement**

```
int    -1  = 0xFFFFFFFFFFFFFFFF
size_t      = 18 446 744 073 709 551 615
```

```
if (uzunluk > 16) return -1; /* -1 bu denetimi GEÇER */
memcpy(hedef, kaynak, uzunluk); /* int -> size_t */
```

Demo 4 — Signed length (CWE-195)


```
ADIM 0  GCC -Wall -Wextra: sessiz · -Wsign-conversion: uyarı  
        MSVC (C): sessiz · aynı kod C++ olarak: C4365  
ADIM 1  kopya 8      -> Kopyalandı: ORNEK-KA  
ADIM 2  kopya 100    -> Reddedildi  
ADIM 3  kopya -1     -> Linux: SIGSEGV (139) · Windows: 0xC0000409  
ADIM 4  ASan         -> negative-size-param: (size=-1)  
ADIM 5  kopya_guvenli -1 / 12abc -> Reddedildi
```

Rule: sizes in `size_t` · **lower and upper** bound · `strtol` · **turn on** warnings

The signed length bug — diagram

İşaretili uzunluk hatası: aynı bitler, iki farklı sayı

int işaretili, size_t işaretsiz — dönüşüm sessizdir





7. Memory management, partitioning, and secure processing

The lifecycle of dynamic memory

Allocate → validate → use → erase secret → free once → forget the pointer

Step	If forgotten	CWE
Validate	<code>malloc</code> returns NULL, unchecked → crash	476
Use	Out-of-bounds write	787
Erase secret	The block goes to another piece of code on the next <code>malloc</code>	226
Free once	Double free	415
Forget the pointer	Use-after-free (UAF)	416
Never free	Leak → denial of service	401

The four lifetimes of memory


Lifetime	How created?	When does it end?	Typical bug
Static	Global variable, <code>static</code> local	When the program ends	Multiple threads writing at once
Automatic	Local variable inside a function	On return	Returning the address of a local variable
Dynamic	<code>malloc</code> / <code>calloc</code> / <code>new</code>	On <code>free</code> / <code>delete</code>	Leak, double free, UAF
Thread	<code>_Thread_local</code> / <code>thread_local</code>	When the thread ends	—

This week's overflow and UAF examples all relate to the **dynamic** lifetime; knowing that alone isn't enough.

The lifetime of memory and its CWEs — diagram

Belleğin dört ömrü ve her adımın CWE'si

her adımın atlanması ayrı bir zafiyet sınıfı üretir



Altı adımın hepsi yapılmazsa bellek hatası kaçınılmazdır.

UAF: why is it so dangerous?

```
Kullanici *aktif = malloc(sizeof *aktif);
Kullanici *onbellek = aktif;           /* ikinci sahip */
free(aktif); aktif = NULL;             /* onbellek hâlâ eski adreste */

char *not = malloc(sizeof(Kullanici)); /* aynı blok geri verilebilir */
strcpy(not, "....");
if (onbellek->yetki) { ... }           /* 'not'un baytlarını okuyor */
```

- `free` doesn't return memory to the OS, it puts it on the **free list** → the same block gets handed out again
-  **Ownership rule:** every block has **one owner**, only that owner frees it

Fixed patterns

```
Kayit *k = calloc(n, sizeof(Kayit));          /* çarpım taşmasını denetler + sıfırlar */  
if (k == NULL) return HATA_BELLEK;  
  
Kayit *yeni = reallocarray(k, n2, sizeof(Kayit));  
if (yeni == NULL) { free(k); return HATA_BELLEK; } /* k'yi kaybetme */  
k = yeni;
```

- ✗ `tampon = realloc(tampon, n);` → on failure, the old block is lost
- ✗ `realloc` on a buffer holding a secret → the moved-away old block **isn't zeroed**

Erasing secrets (Recipe 13.2)

Platform	Function
Linux, BSD	<code>explicit_bzero(p, n)</code>
Windows	<code>SecureZeroMemory(p, n)</code>
C11 Annex K / C23	<code>memset_s</code> / <code>memset_explicit</code>
None available	Byte by byte through a <code>volatile</code> pointer

Hidden copies of a secret: structs passed by value · `realloc`'s old block · `std::string` (SSO, vanishes without erasure) · I/O buffers · swap and hibernation files

Keep it off disk (Recipe 13.3) and RAI

```
mlock(anahtar, sizeof anahtar);           /* Windows: VirtualLock */
/* ... kullan ... */
cen429_sil(anahtar, sizeof anahtar);
munlock(anahtar, sizeof anahtar);
```

```
class GizliAnahtar {                       /* yıkıcı otomatik siler */
    ~GizliAnahtar() { cen429_sil(v_.data(), v_.size()); }
    GizliAnahtar(const GizliAnahtar&) = delete; /* kopya = ikinci sır */
};
```

Locking is at the **page** level, and limited in amount → only a small key region

Memory-bug hunters

Tool	What it finds	Use
AddressSanitizer	Overflow, UAF, double free	<code>-fsanitize=address</code> · <code>/fsanitize=address</code>
LeakSanitizer	Leaks	With ASan (Linux)
MemorySanitizer	Uninitialized reads	Clang <code>-fsanitize=memory</code>
Valgrind	Everything (slow)	<code>valgrind --leak-check=full ./prog</code>
CRT debug heap	Leaks	MSVC <code>_CrtSetDbgFlag</code>

Evaluator: dumps memory after the process ends, searches for the test password

Protected code partitioning: shrink the TCB

Trusted Computing Base (TCB): all the code where "if it's buggy, security collapses"

- A 200,000-line application, a 2,000-line **protected core**
- Review, obfuscation, integrity checking only for the core → cost stays there too

Bölümleme: güvenilir hesaplama tabanını küçült

aralarından yalnız tutamaç ve şifreli veri geçer

GENEL BÖLGE

büyük · hızlı değişir

Arayüz

İş mantığı

hata olasılığı YÜKSEK

KORUNAN ÇEKİRDEK

küçük · seyrek değişir

Dar arayüz (tutamaçlar)

Kripto ve anahtar yönetimi

Bütünlük / RASP denetimleri

incelenebilir, kanıtlanabilir

Güvenlik kritik kod ne kadar küçükse, doğruluğunu kanıtlamak o kadar kolaydır.

Partitioning options

Option	Isolation	Example
Separate module, same process	Low	A native library in a Java app
Separate process	Medium	Browser sandbox, privileged helper process (Recipe 1.4)
OS service	Medium–high	DPAPI, Keychain, Android Keystore
Hardware	High	Secure element, TPM, TEE
Move to server	Highest	The secret never goes to the client, stays in the HSM

Strongest: **remove the secret from the client**; failing that, hardware; failing that, **software protection** (this course's topic)

Interface: narrow, validated, secret-free

```
typedef struct KasaAnahtari KasaAnahtari;    /* opak tür */

int kasa_ac(const char *parola, size_t n, const unsigned char tuz[16],
            KasaAnahtari **cikti);           /* parolayı siler */
int kasa_coz(KasaAnahtari *a, const unsigned char *sifreli, size_t n,
             unsigned char *acik, size_t kapasite, size_t *yazilan);
void kasa_kapat(KasaAnahtari *a);           /* anahtarı siler */
```

1. **Don't hand out the secret, do the work** (a handle)
2. **Validate** everything that crosses the boundary
3. Never pass **plain data** across the boundary

Oracles and code lifting

- **Decryption oracle:** the attacker calls `kasa_coz` with their own data — no key needed
- **Code lifting:** the protected function is copied as-is and run like a black box

Countermeasure	Idea
Bind to device/instance	Wrong result in another environment
Bind to version	Key won't decrypt on an old/modified version
Context checking	The caller's integrity and call order are verified
No shared static key	Breaking one copy doesn't affect others

Secure processing under encryption: the window of exposure

Kötü: [aç]  [kapat] anahtar hep açık
İyi: [aç]  [kapat] yalnız işlem anında

1. **Just-in-time decryption** → erase before the function returns
2. **Masking**: key XOR a random mask — two meaningless pieces in memory
3. **Splitting**: kept in different places, combined only on use
4. **Code encryption**: function code decrypted at runtime (weeks 4, 9)
5. **White-box**: the key is never exposed at all (week 11)

Masking in memory — code example

```
typedef struct {
    uint8_t maske[32];    /* açılışta rastgele üretilir */
    uint8_t ortulu[32];   /* anahtar XOR maske */
} OrtuluAnahtar;


static void anahtari_ac(const OrtuluAnahtar *o, uint8_t gecici[32])
{
    for (size_t i = 0; i < 32; i++)
        gecici[i] = o->ortulu[i] ^ o->maske[i];
}
```

What sits in memory is **not the key itself**, but two meaningless pieces; combined only when used.

Window of exposure — diagram

Açıklık penceresi: sır ne kadar süre açık?

aynı işi yapan iki tasarım, çok farklı risk



Bellek dökümü alan saldırgan için önemli olan tek şey: o anda anahtar açık mıydı?

The secure development lifecycle

Stage	Activity	Week
Requirement	Security requirements from standards	13
Design	Threat model, protection plan	1–2
Implementation	Coding rules, banned functions, static analysis	4–5
Verification	Sanitizer, fuzzing, penetration testing	4, 12
Release	Signed build, unique version identity	1, 10
Response	Vulnerability disclosure, emergency update	12

Microsoft SDL · NIST SSDF (SP 800-218)

Unique version identity

Is what was reviewed the same as what was shipped? → 1.4.2 + git id + the binary's SHA-256

```
execute_process(COMMAND git rev-parse --short=12 HEAD  
                OUTPUT_VARIABLE GIT_KIMLIK OUTPUT_STRIP_TRAILING_WHITESPACE)  
target_compile_definitions(kasa PRIVATE SURUM="1.4.2" GIT_KIMLIK="${GIT_KIMLIK}")
```

```
Get-FileHash .\kasa.exe -Algorithm SHA256 · sha256sum ./kasa
```

In the field, the version also feeds into **key derivation** → a version-rollback attack won't work


Change management: seven steps

1. Baseline → 2. Change request → 3. Classification → 4. Approval and planning → 5. Development and testing → 6. Release → 7. Verification
- Between 3 and 4, a **security impact analysis**: which asset, interface, threat, countermeasure?
 - Reviewing only the changed part = **delta assessment** (week 13)
 - Repository: protected branches · at least one reviewer · signed tag · MFA

Change management — diagram

Değişiklik yönetimi: yedi adım

her değişiklik izlenebilir ve geri alınabilir olmalı



Sertifikasyonda “hangi sürüm onaylandı?” sorusunun cevabı bu döngüdür (12. hafta).

Trade-off log

Decision	Chosen	Reason	Residual risk
Logging in release	Fully removed by a build macro	Silenced code remains as strings	Hard to diagnose in the field
Error codes	Only success/failure (opaque)	Detail gives the attacker a hint	A separate diagnostic channel
Native "debuggable" flag	On (justified exception)	Integrity checking needs to read its own memory	Compensated with debugger detection
Password derivation time	250 ms	Balance between wait time and brute force	A very weak password

An unwritten trade-off is an **error** in an evaluator's eyes

This week's toolbox

Task	Linux / GCC	Windows / MSVC
Suspicious-code warnings	<code>-Wall -Wextra -Wsign-conversion</code>	<code>/W4</code>
Library size checking	<code>-D_FORTIFY_SOURCE=2 -O2</code>	<code>_s</code> functions like <code>strcpy_s</code>
Catching memory bugs	<code>-fsanitize=address</code>	<code>/fsanitize=address</code>
Memory dump	<code>gdb + gcore</code>	<code>MiniDumpWriteDump</code>
Machine code	<code>objdump -d</code>	VS Disassembly
Erasure that can't be removed	<code>explicit_bzero</code>	<code>SecureZeroMemory</code>
Running a program safely	<code>posix_spawn</code> + absolute path	<code>CreateProcessW</code> + full path

This week's toolbox — diagram

Bu haftanın araç kutusu

hangi araç hangi hatayı yakalar

<code>-Wall -Wextra</code>	şüpheli kod uyarısı	derleme
<code>-Wsign-conversion</code>	işaretili/işaretsiz dönüşüm	derleme
<code>-D_FORTIFY_SOURCE=2</code>	boyutu bilinen tampon denetimi	çalışma
<code>-fsanitize=address</code>	sınır dışı, use-after-free	çalışma
<code>gdb / gcore</code>	bellek dökümü alma	inceleme
<code>strings / objdump -d</code>	ikilideki metin ve kod	inceleme

Derleme zamanı ucuz, çalışma zamanı kesin, inceleme araçları saldırganın da elinde.



8. Project and wrap-up

Project: this week's tasks

1. **Team** (individual or up to 4 people) and **topic**
2. GitHub repository + **project plan** (work packages, timeline) → get it approved
3. Security guide draft:
 - **S0** cover and version history · **S2** product overview
 - **S3** architecture and **interface table**
 - **S4** attacker model · **STRIDE** · **attack tree**

Class exercise (15 min)

Student grading system: the instructor enters grades, the student sees them, the data lives on a server.

1. Data flow diagram + trust boundaries
2. **One** threat from each STRIDE letter
3. Attack tree for "raise my grade from 45 to 85" — which is the **cheapest** path?

End to End: Finishing the Vault's Defence

Section 4's threat table for the "Vault" had two open items; the rest of this week closes them with **actual code**:

Vault's threat	This week's answer
T2 — password/key left in memory	Section 5: <code>explicit_bzero</code> / <code>SecureZeroMemory</code> , <code>mlock</code> , disabling core dumps (Demo 2)
T8 — a memory bug in the parser	Sections 6–7: bounds checking, ASan, the safe pattern (Demo 3–4)

Security is not a **feeling**, it's a **process**: know the asset, model the threat, protect it in layers, write down the residual risk. We'll use this framework all term.

Classic Mistakes — Summary

Mistake	Section	Rule
Trusting PATH/the environment (CWE-426)	5	Absolute path + a clean environment + no shell
A secret left in memory (CWE-14)	5	<code>explicit_bzero</code> / <code>SecureZeroMemory</code> , in a way the compiler can't remove
A buffer overflow (CWE-121/787)	6	Validate first, then copy with a bound; avoid dangerous functions
The off-by-one bug (CWE-193)	6	Use <code><</code> , not <code><=</code> ; always leave room for the terminator
A signed/unsigned integer bug (CWE-195)	6	Keep sizes as <code>size_t</code> ; check both a lower and an upper bound
A leak, double free, or UAF (CWE-401/415/416)	7	Give every block a single owner; set the pointer to <code>NULL</code> after <code>free</code>

All six mistakes were covered earlier in this deck with a demo or code example; here they are gathered in one glance.

On Your Own (Not Graded)

1. The PATH trap: repeat with `ls` (Linux) / `whoami` (Windows), fix with a full path
2. In Demo 2, `MOD optimize` → `MOD korumasiz`: does `memset` work now?
3. In Demo 3, 16, 17, 18, 19, 20 characters: a table of `yonetici` values
4. Add `bakiye` to `struct oturum`: push it above 1000 with an overflow, then prevent it
5. A STRIDE table for your own project (≥ 8 threats + countermeasures)
6. Read the Heartbleed report: which check was missing? Which CWE?

Details and hints: course site → Week 1 → "Try it yourself"

Self-check

1. Difference between asset, threat, vulnerability?
2. The recipient IBAN changing in transit — which CIA property?
3. What does a white-box attacker do differently?
4. STRIDE "E" — which demo?
5. AND / OR node — which is good for defence?
6. Why is `system("date")` dangerous? Three fixes?
7. Why did `memset` disappear?
8. Why didn't ASan catch the 17 characters?
9. What happens when `-1` reaches `memcpy` ?
10. Defence in depth — an example from the mobile payment architecture?

Self-check — answers (1–10)

1. **Asset** is a valuable thing to protect; **threat** is a possible harmful event/actor; **vulnerability** is the flaw a threat can exploit.
2. **Integrity** — an IBAN changing in transit is a data integrity violation.
3. **White-box (MATE) owns** the program (reads memory/key, modifies code); a network attacker only connects **from outside**.
4. **E = Elevation of Privilege** → the privilege-dropping / `setuid` demo.
5. **AND** is good for defence: the goal needs all sub-steps, cutting **one** stops the attack. OR is convenient for the attacker.
6. `system` interprets a **shell** → command injection/PATH hijacking. Fix: call **directly** with `execvp` (no shell), a **full path**, an input allowlist.
7. The compiler removed it as a **dead store** (never read afterward). Use `memset_s` / `explicit_bzero` / `volatile`.
8. The overflow didn't **reach** the redzone / that code path wasn't run; ASan only sees the **instrumented** region.
9. `-1` becomes **SIZE_MAX** in a `size_t` → a huge copy → overflow/crash.
10. Payment: TLS+pinning (transit) · AEAD (at rest) · RASP/anti-debug · white-box/HSM key · server checks — **layers**.

Self-check (continued)

1. Why is **S** never asked of a data flow?
2. What are the four responses to a threat? Give an example of "eliminate"
3. Why is the **group dropped first** in a setuid program?
4. Why isn't `strncpy` a safe `strcpy` ?
5. Why did Heartbleed leak data without anything crashing?

Self-check — answers (continued 1–5)

1. **S is not asked** of a data flow: a flow has no identity (spoofing is asked of endpoints/assets).
2. **Mitigate · eliminate · transfer · accept**. "Eliminate" example: **never offering** a "remember my password" feature at all.
3. In `setuid`, the **group is dropped first**; if the user is dropped first, there's **no privilege left** to change the group.
4. `strncpy` **doesn't add** a terminator on a long source → an unterminated string.
5. Heartbleed only **read**; the memory it read was the process's own **valid** memory → it leaked without crashing.

Self-check (continued)

6. What does a canary do when it catches an overflow? Why is it still a problem?
7. What does `free(p); p = NULL;` prevent, and not prevent?
8. What does `mlock` prevent, and not prevent?
9. What is a "decryption oracle"? A countermeasure?
10. The three components of unique version identity?

Self-check — answers (continued 6–10)

6. When a canary is corrupted it **terminates** the program (denial of service); it does **not see** the overflow into the neighboring variable.
7. `free(p); p=NULL;` prevents **double free/UAF**; it does **not** protect against **aliases** (another pointer to the same address).
8. `mlock` prevents memory from being written to **swap**; it does not prevent it from being **read**, or a secret from staying unerased.
9. **Decryption oracle**: the ability to decrypt without the key. Countermeasure: context/integrity checking, a call limit, **binding to the device**.
10. Unique version: **semantic version + git id + binary SHA-256**.

Glossary

Term	Meaning
CIA	Confidentiality/integrity/availability
MATE	An attacker who owns the program
STRIDE	Six threat classes
Attack tree	Breaking a goal into sub-steps
DFD	Data flow diagram
Defence in depth	Many countermeasures together

Next Week

Week 2 — Computer Viruses and Security Models

We will apply the threat-modelling tools we built this week (STRIDE, attack trees) and the attacker model to a concrete class of threat — malware: how a virus spreads, how it hides, and which access-control models (DAC/MAC/RBAC) stop it.

The CWE classification we met for this week's memory-safety bugs will deepen next week with CVE and CVSS.